
Automating Build, Packaging, and Deployment for the Mcommerce Application Using DevSecOps

Authored by :

Mr. MOARRAF YASSINE

Mr. ELOUAZZANI SOUFIAN

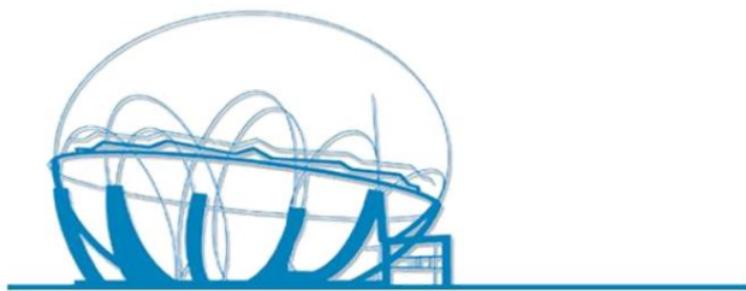
Mr. ELMOUMEN YASSINE

Mr. TOURÉ ABDOULAYE

Mr. TANAFAAT MAROUANE

Under the supervision of:

Prof. Driss ALLAKI



AGENCE NATIONAL DE RÉGLEMENTATION DES TÉLÉCOMMUNICATIONS

INSTITUT NATIONAL DES POSTES ET TÉLÉCOMMUNICATIONS

Class of : 2024/2025

Project repositories:

Main repository: <https://gitlab.com/iccn7501450/cied-demo>

IaC repository: https://gitlab.com/adrian_v/product-microservice-kubernetes

List of Figures

2.1	Initial .gitlab-ci.yml	5
2.2	Compilation and testing job	6
2.3	Compilation and testing job	7
2.4	Private variable.	7
2.5	Private token in DockerHub	8
2.6	CI pipeline configuration inside .gitlab-ci.yml	8
2.7	The CI pipeline running	9
2.8	The CI pipeline running in more detail	9
2.9	The maven-build job in more detail	9
2.10	The docker-build job in more detail	10
3.1	SSH Connection	12
3.2	Upgrades Packages	13
3.3	Adding certificates	13
3.4	Adding GPG Docker Key	14
3.5	Installing Docker	14
3.6	Docker Ps Command	14
3.7	Downloading Minikube binary	15
3.8	Installing Minikube binary	15
3.9	Starting Minikube	15
3.10	Installing kubectl	15
3.11	Listing nodes by kubectl	16
3.12	Creating Namespace	17
3.13	Cloning Microservice repository	17
3.14	Product-deployment YAML file	18

3.15 Product-service YAML file	19
3.16 Applying the product-service and product-deployment files	20
3.17 Checking Pods	20
3.18 Checking Services	20
3.19 Checking Deployments	21
3.20 Checking ReplicaSet	22
3.21 Displaying Resource Details	22
3.22 Accessing Data Hosted by the Product Microservice	23
4.1	25
4.2	25
4.3 Installing Argo CD	25
4.4 Port forwarding to access Argo CD	26
4.5 Ssh tunneling to access the UI locally	26
4.6	26
4.7	27
4.8	27
4.9	27
5.1	29
5.2	29
5.3	29
5.4	30
5.5	31
5.6	31
5.7	32
5.8 gitlab-runner container is running	32
5.9	32
5.10 The local gitlab-runner is connected to GitLab	32
5.11	33
5.12	33
5.13	34
5.14	34
5.15	35

5.16 Different generic jobs	36
5.17 Example of a generic job with default variable that will be overridden	36
5.18 Main CI configuration where the included generic jobs are customized	36
5.19 Defining Roles	37
5.20 Groups	39
5.21 protection des branches	40
5.22 merge request	40
5.23 CI/CD parameter	41
5.24 members	42
5.25 push the file to the root folder of the project	43
5.26 Excerpt from the .gitlab/pre-commit.yml file	50
5.27 Excerpt from the .pre-commit-config.yaml file	51
5.28 Excerpt from the .gitlab-ci.yml file	53
5.29 GitLab pipeline validated after execution of <code>pre-commit-check</code>	54
5.30 <code>pre-commit run -all-files</code> executing various checks (formatting, checkstyle, etc.)	55
5.31 <code>git commit</code> command triggering <code>pre-commit</code> hooks, followed by <code>git push</code>	56
5.32 Example of Maven tests executed before push	57
5.33 Output of the <code>pre-push</code> hook: 11 tests executed, no failures	57
5.34 Example <code>post-commit</code> script (notification and tagging)	58
5.35 Example <code>pre-push</code> script (tests and checks)	59
5.36 Java class <code>test.java</code> generating a false positive	60
5.37 <code>checkstyle.xml</code> file with rule configuration	61
5.38 Running Checkstyle and the error generated	62
5.39 Resolution of the False Positive	62
5.40	64
5.41	64
5.42	65
5.43 configuration for Docker Scout inside a container	66
5.44 docker scout job results	67
5.45 dast test configuration -1-	68
5.46 dast test configuration -2-	68
5.47 pipline for security testing gone Well	69
5.48	70

5.49 Creating a Slack workplace	70
5.50 Creating a channel for vulnerabilities notifications	71
5.51	71
5.52	72
5.53	73
5.54	74
5.55 Stages.yaml updated with the cleanup stage.	74
5.56 Instructions for the cleanup stage.	75
5.57 Extension of the pipeline job file.	75
5.58 Pipeline execution verification.	76
6.1 Add prometheus chart to Helm repo	77
6.2 Installing Prometheus on Kubernetes	77
6.3 Creation of prometheus components : pods ,service and daemon-set	77
6.4 Prometheus ConfigMap	78
6.5 Secrets	78
6.6 Custom Resource Definition CRD	78
6.7 Pods	78
6.8 Prometheus WEB UI	79
6.9 Prometheus Alerts	79
6.10 Prometheus Taragets	79
6.11 Prometheus yaml file for configurations	79
6.12 Graphana Pod	80
6.13 Access the Grafana application in a cluster using Port Forwarding	80
6.14 Graphana login	80
6.15 alerts	81
6.16 CPU	81
6.17 Memory	81
6.18 Network	81
6.19 Disk usage	82
6.20 Prometheus configuration and metrics scraping.	83
6.21 Grafana dashboard visualizing metrics.	83
6.22 Minikube Preparation and Addition of the Stable Helm Repository	84

6.23 Installation of Elasticsearch with Helm	85
6.24 Initial Status of Elasticsearch Pods	85
6.25 Verification of Elasticsearch Services	86
6.26 Status of Elasticsearch Pods After Stabilization	86
6.27 Fluentd YAML Configuration	87
6.28 Verification of Fluentd Pod	88
6.29 Custom Configuration for Kibana	88
6.30 Deployment of Kibana with Helm	89
6.31 Verification of Kibana Pods	89
6.32 Verification of Kibana Service	90
6.33 Configuration of Port-Forward to Access Kibana	90
6.34 Kibana Home Page	91
6.35 Log Visualization in Kibana	92

List of Tables

Listings

5.1	Installing a pre-push hook with pre-commit	56
-----	--	----

Contents

List of Figures	vi
List of Tables	vii
List of Listings	viii
Table of Contents	xii
1 General Project Context	1
2 Creating a CI pipeline	3
2.1 Overview of GitLab CI/CD	3
2.1.1 jobs	3
2.1.2 Pipelines	3
2.1.3 Runners	4
2.2 Creating a CI pipeline in GitLab	4
2.2.1 The ".gitlab-ci.yml" file	4
2.2.2 Syntax of ".gitlab-ci.yml" file	4
2.2.3 Compilation and testing stage	6
2.2.4 Containerization stage	6
2.3 Testing the result	8
3 Local Deployment with Kubernetes	11
3.1 Overview of Technologies	11
3.2 Installing Dependencies and Creating the Minikube Cluster	12
3.3 Configuring the Namespace	17
3.4 Adding YAML Files for Deployment and Service	17
3.4.1 Deployment YAML File	18

3.4.2	Service YAML File	19
3.4.3	Applying Deployment and Service to the Cluster	20
3.5	Verifying Kubernetes Resources	20
3.5.1	Check Pods:	20
3.5.2	Check Service:	20
3.5.3	Check Deployments:	21
3.5.4	Check ReplicaSet:	22
3.5.5	Check All Components:	22
4	Extension of the CI pipeline with GitOps	24
4.0.1	What is GitOps	24
4.0.2	Using ArgoCD	24
4.0.3	Preparing the IaC repository	24
4.0.4	Argo CD configuration	25
5	Optimisation et Sécurisation du VCS et du CI Pipeline	28
5.1	Optimizing the CI pipeline	28
5.1.1	CI in Local GitLab Runner	28
5.1.2	Job rules	32
5.1.3	Caching	35
5.1.4	Generic jobs	35
5.2	gestion des groupes et permissions	37
5.2.1	Définir des rôles	37
5.2.2	Group definitions with roles	38
5.2.3	Project access configuration	39
5.2.4	Protection des branches	39
5.2.5	Configuration des merge requests	40
5.2.6	CI/CD	41
5.2.7	Configuration des membres	42
5.3	Protected Branches	43
5.3.1	Protected Branch Controls	43
5.3.2	Who Can Modify a Protected Branch	43
5.3.3	Steps to Protect a Branch for All Projects in a Group	44

5.4	Add a Coverage Check Approval Rule	45
5.4.1	Steps to Add a Coverage-Check Approval Rule	45
5.5	Pre-Commit Hooks	49
5.5.1	Integration of Pre-Commit Hooks in GitLab	49
5.5.2	Structure of the <code>pre-commit.yml</code> File in GitLab	49
5.5.3	Configuration of <code>.pre-commit-config.yaml</code>	50
5.5.4	Inclusion in <code>.gitlab-ci.yml</code> and Resolution of the Issue	52
5.5.5	Evaluation of Using Frameworks like <code>pre-commit</code> (Tests for <i>pre-push</i> , <i>post-commit</i>)	55
5.5.5.1	Installing a <code>pre-push</code> Hook via <code>pre-commit</code>	56
5.5.6	Example Scripts for <code>post-commit</code> and <code>pre-push</code>	58
5.5.7	Strategies to Minimize False Positives	60
5.5.8	Step 1: Creating a Java Class with a False Positive	60
5.5.9	Step 2: Configuring Checkstyle Rules	61
5.5.10	Step 3: Running Checkstyle on the Class	62
5.5.11	Step 4: Resolving the False Positive	62
5.5.12	Step 5: Validating the Changes	62
5.6	Vulnerability scanning	63
5.6.1	Static Application Security Testing	64
5.6.2	Software Composition Analysis	64
5.6.3	Container Security	65
5.6.4	Dynamic Application Security Testing	67
5.7	Vulnerability management and notification	69
5.7.1	Vulnerability management	69
5.7.2	Configuring notification using Slack	70
5.8	Automated Pipeline Cleaning	74
6	Observability and Monitoring	77
6.1	infrastructure Monitoring	77
6.1.1	Installing Prometheus on Kubernetes	77
6.1.2	Access Prometheus UI	78
6.1.3	Graphana	80
6.2	Implementation of Application Monitoring for the Produit Microservice	82
6.2.1	Why Prometheus and Grafana?	82

6.2.2	Implementation Steps	82
6.2.2.1	Prometheus Setup	82
6.2.2.2	Grafana Setup	82
6.2.3	Screenshots and Observations	83
6.2.4	Conclusion	83
6.3	Adding Log Monitoring to Analyze Microservice Pod Logs (Using an EFK stack)	84
6.3.1	Preparation of Minikube and Addition of the Stable Helm Repository	84
6.3.2	Deployment of Elasticsearch with Helm	84
6.3.2.1	Installation and Verification of Elasticsearch Pods	85
6.3.2.2	Verification of Elasticsearch Services	85
6.3.2.3	Confirmation of Stability of Elasticsearch Pods	86
6.3.3	Configuration and Deployment of Fluentd as a DaemonSet	86
6.3.3.1	Creation of the Fluentd YAML File	86
6.3.3.2	Verification of Fluentd Pods	87
6.3.4	Deployment of Kibana with a Custom Configuration File	88
6.3.4.1	Configuration of Kibana	88
6.3.4.2	Deployment and Verification of Kibana Pods	88
6.3.4.3	Verification of Service and Access to Kibana	89
6.3.4.4	Access to Kibana Dashboard	90
6.3.4.5	Log Visualization	92
6.3.5	Conclusion	92

Chapter 1

General Project Context

In today's fast-paced and highly competitive software industry, the ability to deliver high-quality solutions quickly and securely is a critical factor for success. Companies are increasingly adopting modern software development practices, such as DevOps and DevSecOps, to streamline their processes, enhance collaboration, and ensure the security and reliability of their applications. For KATA, a young and ambitious software development company, this shift represents both an opportunity and a challenge. While KATA is known for delivering high-quality software solutions, its current delivery process is slower compared to industry competitors, which limits its ability to seize new market opportunities effectively.

To address these challenges, KATA has decided to explore the DevSecOps approach, which integrates development, security, and operations into a seamless workflow. This approach aims to automate and secure the software delivery pipeline, enabling faster and more reliable deployments without compromising on quality. As part of this initiative, KATA has launched a project to develop a Proof of Concept (PoC) for automating the build, packaging, and deployment of a microservices-based application called Mcommerce. This application, composed of three core microservices (Product, Order, and Payment), a web client, and an API Gateway, serves as a testbed for evaluating the effectiveness of DevSecOps practices.

The primary goal of this project is to demonstrate how automation, containerization, and security integration can improve KATA's software delivery process. By implementing a CI/CD pipeline, containerizing the application using Docker, orchestrating deployments with Kubernetes, and integrating security tools into the workflow, the project aims to achieve faster delivery cycles, enhanced security, and greater operational efficiency. The success of this PoC will provide KATA with valuable insights into the potential of DevSecOps and guide its future strategy for adopting modern software development practices.

This report outlines the methodology, implementation, and results of the PoC, highlighting the key achievements, challenges, and lessons learned throughout the project.

Chapter 2

Creating a CI pipeline

2.1 Overview of GitLab CI/CD

CI/CD is a continuous method of software development, where you continuously build, test, deploy, and monitor iterative code changes. This iterative process helps reduce the chance that you develop new code based on failed or buggy previous versions. GitLab CI/CD can catch bugs early in the development cycle and help ensure that the code deployed to production complies with your established code standards. Gitlab provides what are called Gitlab runners processes that pick up and execute CI/CD jobs.

For this part, we will focus only on the CI part of the pipeline to build, test, and containerize the microservice.

2.1.1 jobs

Jobs are a list of commands to execute to accomplish tasks like building, testing, or deploying code. they are independent of other jobs and can be run in parallel. Jobs can have many status like: failed, warning, pending, running.

2.1.2 Pipelines

Pipelines are composed of:

- Global YAML keywords that control the overall behavior of the project's pipelines
- Jobs that execute commands to accomplish a task. For example, a job could compile, test, or deploy code. Jobs run independently from each other, and are executed by runners.

- Stages, which define how to group jobs together. Stages run in sequence, while the jobs in a stage run in parallel. For example, an early stage could have jobs that lint and compile code, while later stages could have jobs that test and deploy code. If all jobs in a stage succeed, the pipeline moves on to the next stage. If any job in a stage fails, the next stage is not (usually) executed and the pipeline ends early.

2.1.3 Runners

Runners are programs that their responsibility is to execute CI/CD jobs. In GitLab there are two types of runners, one's that run on physical machines and other in virtual instances. More concretely, in the ".gitlab-ci.yml" file that we will see next, you can specify a container image that you want to use when running the job. The runner loads the image, clones your project, and runs the job locally or in the container.

2.2 Creating a CI pipeline in GitLab

In this part we will discuss how to create a CI/CD pipeline in Gitlab.

2.2.1 The ".gitlab-ci.yml" file

The file ".gitlab-ci.yml" is a special yaml file that exists in the root directory of a repo. Its the main file that defines a pipeline as it specifies the stages, jobs, and scripts to be executed during the CI/CD pipeline.

2.2.2 Syntax of ".gitlab-ci.yml" file

Generally the file follows the following format

```
stages:
  - build
  - test
  - deploy

variables:
  PROJECT_NAME: "my_project"
```

```

BUILD_DIR: "build"

build:
  stage: build # Part of the build stage
  script: ...

test:
  stage: test # Part of the test stage
  script:
    - echo "Running tests..."
    - echo "Tests passed."

deploy:
  stage: deploy # Part of the deploy stage
  script: ...

```

First we define our stages and then for each stage we defines jobs we want to be executed part of that stage. A typical pipeline have 3 stages: build, test and deploy.

In our case we have a pipeline with two stages: Build/Test and Containerize, thus we will use the following format:

```

1  stages:
2  | - build
3  | - container
4
5
6  maven-build:
7  | stage: build
8
9
10 docker-build:
11 | stage: container
12
13

```

Figure 2.1: Initial .gitlab-ci.yml

2.2.3 Compilation and testing stage

Our microservice is written using Java 18 and uses Maven as a build system and for testing. These factors result in the following build job description.

```

1 stages:
2   - build
3   - container
4
5
6 maven-build:
7   stage: build
8   image: maven:3.8.7-openjdk-18-slim
9   script:
10    - chmod +x ./mvnw
11    - ./mvnw clean package
12   artifacts:
13     paths:
14       - ./target/*.jar
15

```

Figure 2.2: Compilation and testing job

Explication:

- Line 6: Specify our job name.
- Line 7: Specify that this job is part of build stage.
- Line 8: We want our job to be executed in a container that provides JDK version 18 with maven.
- Line 9: Here we define what our job does using shell commands.
- Line 10: This command make our maven binary file executable.
- Line 11: Start the test and build process
- Line 12-13-14: Specify that the content of the target folder is an artifact so it won't be deleted after the job and we can access it in other jobs.

2.2.4 Containerization stage

We want to user docker to containerize our app and Dockerhub as a container registry to share our docker image. for this we will use the following job.

```

16 docker-build:
17   stage: container
18   image: docker
19   services:
20     - docker:dind
21   before_script:
22     - docker login -u projdevopsiccn2 -p $DOCKER_TOKEN
23   script:
24     - docker version
25     - docker build -t projdevopsiccn2/mcommerce .
26     - docker push projdevopsiccn2/mcommerce
27

```

Figure 2.3: Compilation and testing job

Explanation:

- Line 16: Specify our job name.
- Line 17: Specify that this job is part of the container stage.
- Line 18: We want our job to be executed in a container that provides Docker (Nested container situation)
- Line 19: Services are additions to the image we want to use
- Line 20: dind (docker inside docker) specify that we want to have the ability to run nested docker containers.
- Line 21: These are command that get ran before the main script.
- Line 22: Because we want to upload our result image to Dockerhub we need first to login.

The login here uses a private variable DOCKER_TOKEN that we defined in the settings under:
Settings > CI/CD > Variables

The screenshot shows the 'Variables' section of a GitLab project's CI/CD settings. It includes a descriptive text about variables, a note on security, and a table for managing variables. The table has columns for Key, Value, Environments, and Actions. A single variable, DOCKER_TOKEN, is listed with its value masked and expanded.

CI/CD Variables </> 1		Reveal values	Add variable
Key ↑	Value	Environments	Actions
DOCKER_TOKEN	All (default)		🔗 trash
<small>Masked Hidden Expanded</small>			

Figure 2.4: Private variable.

The token can be created from DockerHub under: Settings > Personal access tokens

Description	Scope	Status	Source	Created	Last used	Expiration date	⋮
Gitlab	Read, Write, Delete	Active	Manual	Dec 12, 2024 at 23:30:16	Dec 21, 2024 at 17:25:27	Jan 11, 2025 at 23:59	

Rows per page: 10 ▾ 1-1 of 1 ⏪ ⏩

Figure 2.5: Private token in DockerHub

- Line 23: Here we define what our job does using shell commands
- Line 24: First we check if docker is working by showing its version
- Line 25: Then we build our image using Dockerfile in the repo
- Line 26: Then we push the resulting image to our Dockerhub repository.

2.3 Testing the result

The resulting ".gitlab-ci.yml" file is:

```

1 stages:
2   - build
3   - container
4
5
6 maven-build:
7   stage: build
8   image: maven:3.8.7-openjdk-18-slim
9   script:
10    - chmod +x ./mvnw
11    - ./mvnw clean package
12   artifacts:
13     paths:
14       - ./target/*.jar
15
16 docker-build:
17   stage: container
18   image: docker
19   services:
20     - docker:dind
21   before_script:
22     - docker login -u projdevopsiccn2 -p $DOCKER_TOKEN
23   script:
24     - docker version
25     - docker build -t projdevopsiccn2/mcommerce .
26     - docker push projdevopsiccn2/mcommerce
27

```

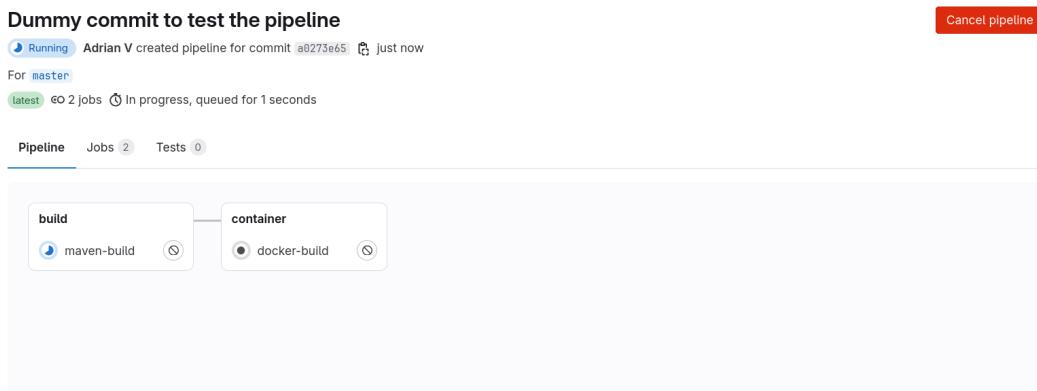
Figure 2.6: CI pipeline configuration inside .gitlab-ci.yml

Now We submit a dummy commit to test our CI pipeline. and the result is :

Figure 2.7 shows our pipeline running, and the commit that triggered it, also the different stages that exists in the pipeline. Figure 2.8 shows in more detail the status of each stage. Figures 2.9 and 2.10 shows log of the jobs "maven-build" and "docker-build" respectively, it also shows the duration, status and other information about each job.

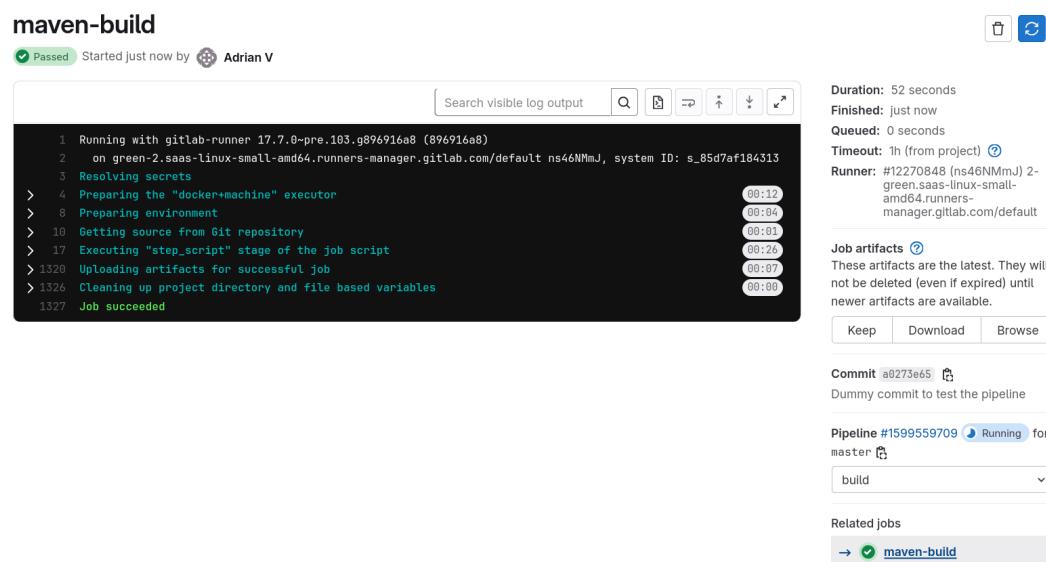
Status	Pipeline	Created by	Stages
Running	Dummy commit to test the pipeline #1599559709 master a0273e65 (latest)		Running Waiting

Figure 2.7: The CI pipeline running



The screenshot shows the GitLab CI pipeline interface. At the top, it displays the pipeline name: "Dummy commit to test the pipeline". Below that, it shows the status: "Running" (Adrian V created pipeline for commit a0273e65 just now). It indicates there are 2 jobs in the master branch, both of which are currently in progress and queued for 1 second. The pipeline structure is shown as a tree: "build" has a child "container" which contains "maven-build" and "docker-build". The "maven-build" job is currently in progress.

Figure 2.8: The CI pipeline running in more detail



This screenshot shows the detailed view of the "maven-build" job. The job status is "Passed" and it started "just now" by Adrian V. The log output shows the execution of the job, including steps like "Running with gitlab-runner", "Preparing the 'docker+machine' executor", and "Job succeeded". The log is timestamped with 1327 entries. To the right of the log, various details are provided: Duration (52 seconds), Finished (just now), Queued (0 seconds), Timeout (1h from project), Runner (a specific runner configuration), and Job artifacts (which are latest and not deletable). There are buttons for Keep, Download, and Browse artifacts. Below the log, it shows the Commit (a0273e65), Pipeline (1599559709), and Related jobs (maven-build).

Figure 2.9: The maven-build job in more detail

The screenshot shows a detailed view of a GitLab CI job named 'docker-build'. The job has passed and was started just now by a user named Adrian V. The log output is displayed in a terminal-like interface with various status indicators and timestamps. Key steps shown include running with a specific gitlab-runner version, preparing the docker+machine executor, resolving secrets, and executing a step script. The duration of the job was 57 seconds. The runner used was #12270857 (ntHFEtyX) 4-green.saas-linux-small-amd64.runners-manager.gitlab.com/default. A dummy commit is mentioned, and the pipeline is #1599559709, which also passed. Related jobs show a dependency on another 'docker-build' job.

```
1 Running with gitlab-runner 17.7.0-pre.103.g896916a8 (896916a8)
2 on green-4.saas-linux-small-amd64.runners-manager.gitlab.com/default ntHFEtyX, system ID: s_8990de21c550
3 Resolving secrets
> 4 Preparing the "docker+machine" executor
> 12 Preparing environment
> 14 Getting source from Git repository
> 21 Downloading artifacts
> 24 Executing "step_script" stage of the job script
> 139 Cleaning up project directory and file based variables
140 Job succeeded
```

Duration: 57 seconds
Finished: just now
Queued: 0 seconds
Timeout: 1h (from project) [?](#)
Runner: #12270857 (ntHFEtyX) 4-green.saas-linux-small-amd64.runners-manager.gitlab.com/default
Commit a0273e65 [🔗](#)
Dummy commit to test the pipeline
Pipeline #1599559709 [Passed](#) for master [🔗](#)
container

Related jobs
→ [Passed](#) docker-build

Figure 2.10: The docker-build job in more detail

Chapter 3

Local Deployment with Kubernetes

3.1 Overview of Technologies

In this chapter, we will explore the deployment of the **Product microservice** using container orchestration and automation tools. also we will use a VPS (Ubuntu machine) so we simulate a real word scenario and different members could connect and work on the same machine .The primary technologies utilized include **Kubernetes**, **Minikube**, and **kubectl**. Below is a brief description of these technologies:

- **Kubernetes**: A powerful open-source container orchestration platform designed to manage the deployment, scaling, and operation of containerized applications. Kubernetes simplifies the management of complex application architectures by automating deployment and ensuring high availability.
- **Minikube**: A lightweight Kubernetes implementation designed for local environments. It allows developers to create a single-node Kubernetes cluster on their machines for development and testing purposes.
- **kubectl**: A command-line tool used to interact with Kubernetes clusters. It provides commands to deploy applications, inspect resources, and troubleshoot issues within a Kubernetes environment.

By leveraging these tools, we demonstrate how to create a local Kubernetes cluster, configure a dedicated namespace, and deploy a microservice using declarative YAML configurations. This hands-on approach not only ensures efficient resource management but also prepares for scalable and production-ready deployments.

3.2 Installing Dependencies and Creating the Minikube Cluster

To begin, we will create a Kubernetes cluster on the VPS (Virtual Private Server) using Minikube. First, we connect to the VPS machine using SSH.

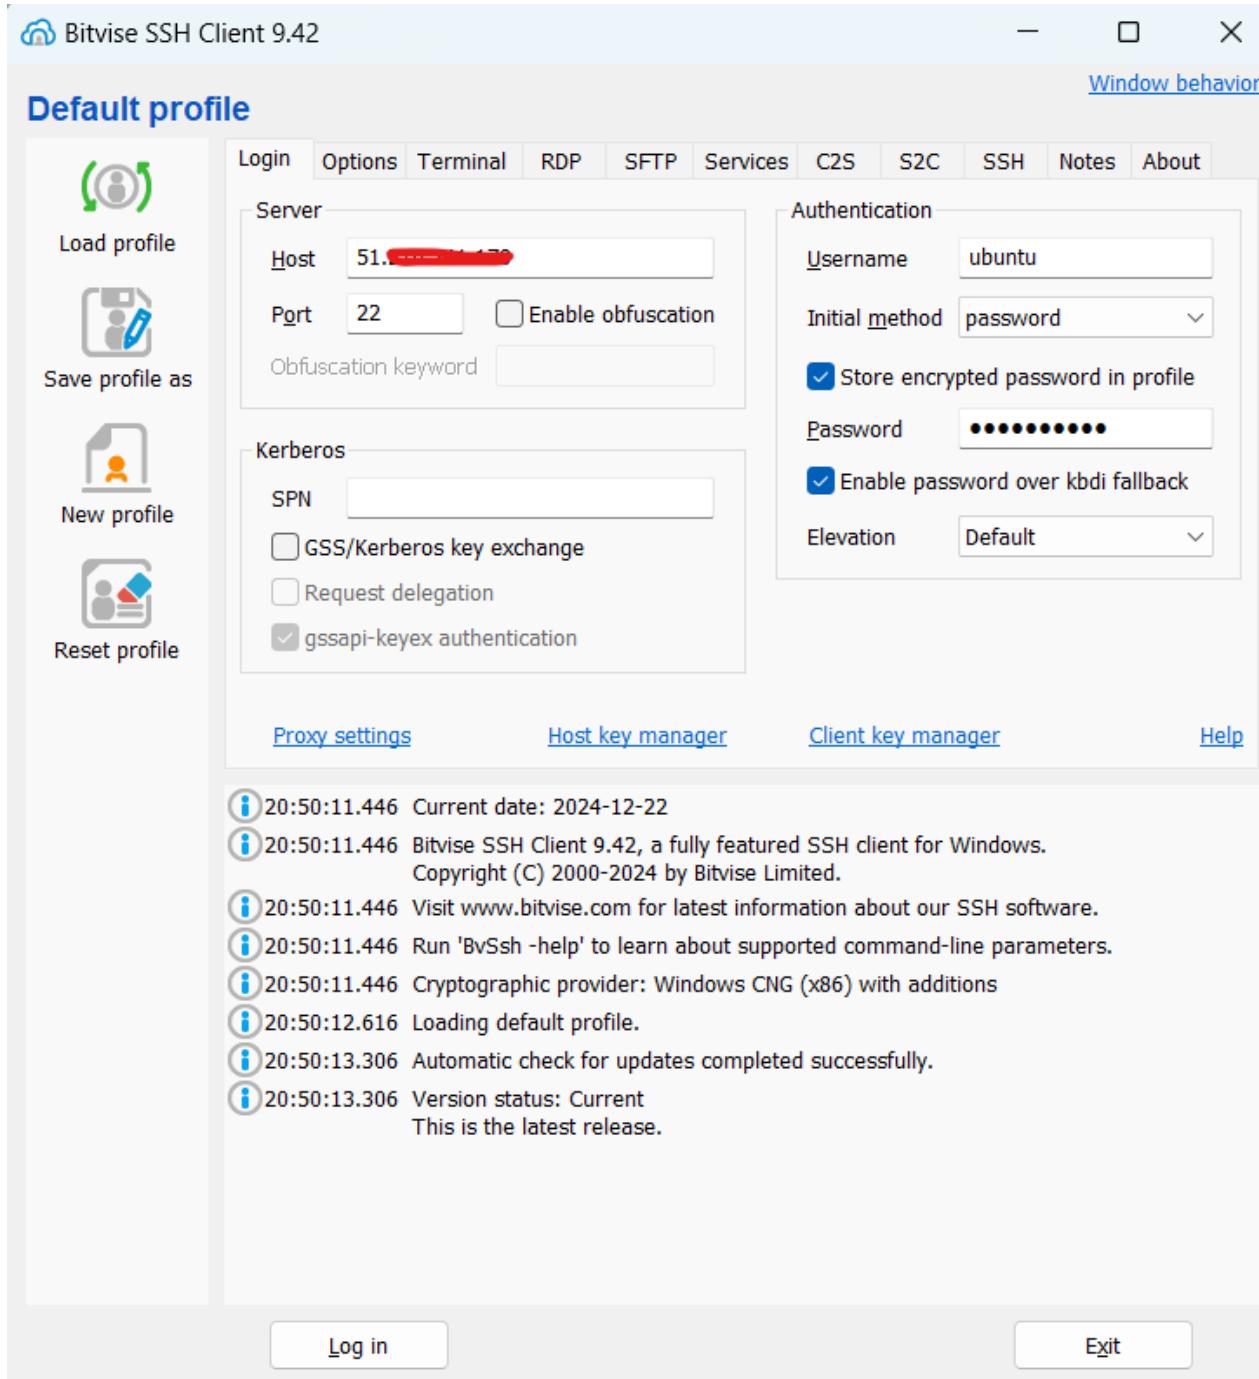


Figure 3.1: SSH Connection

After connecting to the VPS, and before installing Docker, it's a good idea to update the local packages to ensure we have the latest available versions.

```
ubuntu@vps-4f4e3327:~$ sudo apt update && sudo apt upgrade -y
Get:1 http://security.ubuntu.com/ubuntu jammy-security InRelease [129 kB]
Get:2 http://nova.clouds.archive.ubuntu.com/ubuntu jammy InRelease [270 kB]
Get:3 http://security.ubuntu.com/ubuntu jammy-security/main amd64 Packages [1997 kB]
Get:4 http://security.ubuntu.com/ubuntu jammy-security/main Translation-en [317 kB]
Get:5 http://security.ubuntu.com/ubuntu jammy-security/restricted amd64 Packages [2701 kB]
Get:6 http://security.ubuntu.com/ubuntu jammy-security/restricted Translation-en [470 kB]
Get:7 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-updates InRelease [128 kB]
Get:8 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-backports InRelease [127 kB]
Get:9 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-updates/main amd64 Packages [2241 kB]
Get:10 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-updates/main Translation-en [378 kB]
Get:11 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-updates/restricted amd64 Packages [2817 kB]
]
Get:12 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-updates/restricted Translation-en [490 kB]
Get:13 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-backports/main amd64 Packages [91.5 kB]
Get:14 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-backports/universe amd64 Packages [31.4 kB]
```

Figure 3.2: Upgrades Packages

Next, we will install the required prerequisite packages.

```
ubuntu@vps-4f4e3327:~$ sudo apt-get install -y apt-transport-https ca-certificates curl gpg
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
ca-certificates is already the newest version (20240203~22.04.1).
ca-certificates set to manually installed.
curl is already the newest version (7.81.0-1ubuntu1.20).
curl set to manually installed.
gpg is already the newest version (2.2.27-3ubuntu2.1).
gpg set to manually installed.
The following packages were automatically installed and are no longer required:
  linux-headers-5.15.0-127 linux-headers-5.15.0-127-generic linux-image-5.15.0-127-generic
    linux-modules-5.15.0-127-generic
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  apt-transport-https
```

Figure 3.3: Adding certificates

Now, let's add Docker's GPG key and the Docker repository to our system.

```
ubuntu@vps-4f4e3327:~$ sudo install -m 0755 -d /etc/apt/keyrings
ubuntu@vps-4f4e3327:~$ sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o /etc/apt/keyrings/docker.asc
ubuntu@vps-4f4e3327:~$ sudo chmod a+r /etc/apt/keyrings/docker.asc
ubuntu@vps-4f4e3327:~$ echo \
> "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.asc] https://download.docker.com/linux/ubuntu \
> $(. /etc/os-release && echo "$VERSION_CODENAME") stable" | \
> sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
ubuntu@vps-4f4e3327:~$ sudo apt-get update
Hit:1 http://security.ubuntu.com/ubuntu jammy-security InRelease
Hit:2 http://nova.clouds.archive.ubuntu.com/ubuntu jammy InRelease
Get:3 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-updates InRelease [128 kB]
Get:4 https://download.docker.com/linux/ubuntu jammy InRelease [48.8 kB]
Get:5 http://nova.clouds.archive.ubuntu.com/ubuntu jammy-backports InRelease [127 kB]
Get:6 https://download.docker.com/linux/ubuntu jammy/stable amd64 Packages [42.5 kB]
Fetched 347 kB in 1s (475 kB/s)
Reading package lists... Done
```

Figure 3.4: Adding GPG Docker Key

```
ubuntu@vps-4f4e3327:~$ sudo apt-get install docker-ce docker-ce-cli containerd.io docker-buildx-plugin
docker-compose-plugin
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  docker-ce-rootless-extras libltdl7 libslirp0 pigz slirp4netns
Suggested packages:
  aufs-tools cgroupfs-mount | cgroup-lite
The following NEW packages will be installed:
  containerd.io docker-buildx-plugin docker-ce docker-ce-cli docker-ce-rootless-extras
  docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
0 upgraded, 10 newly installed, 0 to remove and 1 not upgraded.
Need to get 124 MB of archives.
After this operation, 446 MB of additional disk space will be used.
Do you want to continue? [Y/n] Y
Get:1 https://download.docker.com/linux/ubuntu jammy/stable amd64 containerd.io amd64 1.7.24-1 [29.5 MB]
Get:2 http://nova.clouds.archive.ubuntu.com/ubuntu jammy/universe amd64 pigz amd64 2.6-1 [63.6 kB]
Get:3 http://nova.clouds.archive.ubuntu.com/ubuntu jammy/universe amd64 libltdl7 amd64 2.4.6-1Build2 [70.6 kB]
```

Figure 3.5: Installing Docker

We then verify that Docker is successfully installed:

```
ubuntu@vps-4f4e3327:~$ sudo docker ps
CONTAINER ID   IMAGE      COMMAND     CREATED      STATUS      PORTS      NAMES
ubuntu@vps-4f4e3327:~$ sudo docker ps -a
CONTAINER ID   IMAGE      COMMAND     CREATED      STATUS      PORTS      NAMES
89ad04b7fd33   hello-world  "/hello"    47 seconds ago  Exited (0) 46 seconds ago
                                         heuristicfaraday
```

Figure 3.6: Docker Ps Command

Now, let's download the latest Minikube binary:

```
zenlnex@vps-4f4e3327:/home/zenlnex$ curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
% Total    % Received % Xferd  Average Speed   Time     Time     Time  Current
          Dload  Upload Total   Spent   Left  Speed
100 99.0M  100 99.0M    0      0  26.2M      0  0:00:03  0:00:03 --:--:-- 26.2M

```

Figure 3.7: Downloading Minikube binary

Once the binary is downloaded, we can install Minikube:

```
zenlnex@vps-4f4e3327:/home/zenlnex$ sudo install minikube-linux-amd64 /usr/local/bin/minikube
zenlnex@vps-4f4e3327:/home/zenlnex$ minikube version
minikube version: v1.34.0
commit: 210b148df93a80eb872ecbeb7e35281b3c582c61
```

Figure 3.8: Installing Minikube binary

Next, let's create the Kubernetes cluster locally using Minikube:

```
zenlnex@vps-4f4e3327:/home/zenlnex$ sudo minikube start --driver=docker --force
* minikube v1.34.0 on Ubuntu 22.04 (kvm/amd64)
! minikube skips various validations when --force is supplied; this may lead to unexpected behavior
* Using the docker driver based on user configuration
* The "docker" driver should not be used with root privileges. If you wish to continue as root, use --force.
* If you are running minikube within a VM, consider using --driver=none:
*   https://minikube.sigs.k8s.io/docs/reference/drivers/none/
* Using Docker driver with root privileges
* Starting "minikube" primary control-plane node in "minikube" cluster
* Pulling base image v0.0.45 ...
* Downloading Kubernetes v1.31.0 preload ...
  > preloaded-images-k8s-v18-v1...: 326.69 MiB / 326.69 MiB  100.00% 24.57 M^[[B
  > gcr.io/k8s-minikube/kicbase...: 486.80 MiB / 487.90 MiB  99.77% 32.81 Mi
* Creating docker container (CPUs=2, Memory=2200MB) ...
* Preparing Kubernetes v1.31.0 on Docker 27.2.0 ...
  - Generating certificates and keys ...
  - Booting up control plane ...
  - Configuring RBAC rules ...
* Configuring bridge CNI (Container Networking Interface) ...
* Verifying Kubernetes components...
  - Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: storage-provisioner, default-storageclass
* kubectl not found. If you need it, try: 'minikube kubectl -- get pods -A'
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
```

Figure 3.9: Starting Minikube

After installing the latest version of kubectl using Snap:

```
zenlnex@vps-4f4e3327:/home/zenlnex$ sudo snap install kubectl --classic
kubectl 1.31.4 from Canonical✓ installed
zenlnex@vps-4f4e3327:/home/zenlnex$
```

Figure 3.10: Installing kubectl

Finally, let's check that kubectl is installed successfully:

```
zenlnex@vps-4f4e3327:/home/zenlnex$ sudo kubectl get nodes
NAME        STATUS    ROLES          AGE    VERSION
minikube   Ready     control-plane  11m    v1.31.0
```

Figure 3.11: Listing nodes by kubectl

3.3 Configuring the Namespace

Let's configure a namespace named `product-local-use` to logically separate the resources related to the Product microservice. using the command:

```
zenlnex@vps-4f4e3327:/home/zenlnex$ sudo kubectl create namespace product-local-use
namespace/product-local-use created
```

Figure 3.12: Creating Namespace

The namespace provides a dedicated space to manage all Kubernetes objects for the microservice.

3.4 Adding YAML Files for Deployment and Service

The deployment and service YAML files manage the microservice's lifecycle and expose it to external access:

- **Deployment YAML:** Configures the number of replicas, container image, and resource specifications for the Product microservice.
- **Service YAML:** Defines the network access point for the microservice using a ClusterIP.

Now, let's clone our `cicd-demo` repository from GitLab in order to add the YAML files for the deployment.

```
ubuntu@vps-4f4e3327:/mcommerce$ sudo git clone https://gitlab.com/iccn7501450/cicd-demo.git
Cloning into 'cicd-demo'...
Username for 'https://gitlab.com': g1d3on_wh1t3
Password for 'https://g1d3on_wh1t3@gitlab.com':
remote: Enumerating objects: 130, done.
remote: Counting objects: 100% (60/60), done.
remote: Compressing objects: 100% (58/58), done.
remote: Total 130 (delta 31), reused 0 (delta 0), pack-reused 70 (from 1)
Receiving objects: 100% (130/130), 23.37 KiB | 1.80 MiB/s, done.
Resolving deltas: 100% (49/49), done.
ubuntu@vps-4f4e3327:/mcommerce$ ls
cicd-demo
```

Figure 3.13: Cloning Microservice repository

After that, let's create a `.Kubernetes` folder that will contain the `product-deployment` and `product-service` YAML files for deployment and service.

3.4.1 Deployment YAML File

The following YAML file describes the deployment configuration for the **mcommerce** microservice in a Kubernetes cluster:

```
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ cat product-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mcommerce
  labels:
    app: mcommerce
spec:
  replicas: 2
  selector:
    matchLabels:
      app: mcommerce
  template:
    metadata:
      labels:
        app: mcommerce
    spec:
      containers:
        - name: mcommerce
          image: projdevopsiccn2/mcommerce:latest
          ports:
            - containerPort: 9001
```

Figure 3.14: Product-deployment YAML file

Now, let's describe this file:

- **apiVersion** and **kind: Deployment**: Specifies that this YAML file defines a Deployment, which ensures the desired number of application instances are running.
- **metadata**: Contains identifying information, including the deployment name **mcommerce** and a label.
- **spec.replicas: 2**: Ensures that two pod replicas of the **mcommerce** microservice are running for redundancy and load balancing.
- **spec.selector.matchLabels**: Matches pods labeled with `app: mcommerce`, ensuring that the deployment manages the intended pods.
- **spec.template**: Defines the pod template, including:

- **metadata.labels**: Assigns the label `app: mcommerce` to the pods created by this deployment.
- **spec.containers**: Configures the container, specifying:
 - * **name: mcommerce**: The container name.
 - * **image: projdevopsiccn2/mcommerce:latest**: The container image to use.
 - * **ports.containerPort: 9001**: Exposes port 9001 inside the container.

This configuration ensures that the `mcommerce` microservice is deployed with two replicas and is accessible on port 9001.

3.4.2 Service YAML File

The following YAML file describes the service configuration for exposing the `mcommerce` microservice in a Kubernetes cluster:

```
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ cat product-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: product-service
spec:
  selector:
    app: mcommerce
  ports:
  - protocol: TCP
    port: 80
    targetPort: 9001
```

Figure 3.15: Product-service YAML file

Now, let's describe this file:

- **apiVersion: v1** and **kind: Service**: Specifies that this YAML defines a Service, which provides a stable network interface to access the microservice.
- **metadata**: Contains the service name (`product-service`), which identifies this resource.
- **spec.selector**: Matches pods labeled with `app: mcommerce`, ensuring that the service routes traffic to the appropriate pods managed by the deployment.
- **spec.ports**: Configures the ports for communication:
 - **protocol: TCP**: Specifies TCP as the communication protocol.

- **port: 80:** Defines the external port for accessing the service.
- **targetPort: 9001:** Maps traffic from the external port 80 to the container's internal port 9001.

3.4.3 Applying Deployment and Service to the Cluster

Let's apply our YAML files to the cluster using the following command:

```
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ kubectl apply -f product-deployment.yaml -n product-local-use
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ sudo kubectl apply -f product-service.yaml -n product-local-use
deployment.apps/mcommerce created
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ sudo kubectl apply -f product-service.yaml -n product-local-use
service/product-service unchanged
```

Figure 3.16: Applying the product-service and product-deployment files

3.5 Verifying Kubernetes Resources

To confirm the successful creation of resources, we use the following kubectl commands:

3.5.1 Check Pods:

```
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ sudo kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
mcommerce-74d49b66fc-cbplx   1/1     Running   0          2m8s
mcommerce-74d49b66fc-g6bmb   1/1     Running   0          2m8s
```

Figure 3.17: Checking Pods

The screenshot demonstrates the pods running as expected.

3.5.2 Check Service:

```
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ sudo kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1      <none>        443/TCP      77m
product-service  ClusterIP  10.98.190.212  <none>        80/TCP      3m58s
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$
```

Figure 3.18: Checking Services

The screenshot highlights the service exposing the microservice.

3.5.3 Check Deployments:

```
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ sudo kubectl get deployments
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mcommerce  2/2     2           2           8m31s
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ sudo kubectl describe deployments mcommerce
Name:           mcommerce
Namespace:      default
CreationTimestamp:  Wed, 18 Dec 2024 15:53:31 +0000
Labels:          app=mcommerce
Annotations:    deployment.kubernetes.io/revision: 1
Selector:        app=mcommerce
Replicas:       2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=mcommerce
  Containers:
    mcommerce:
      Image:      projdevopsiccn2/mcommerce:latest
      Port:       9001/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
      Node-Selectors: <none>
      Tolerations:  <none>
  Conditions:
    Type      Status  Reason
    ----      -----  -----
    Available  True    MinimumReplicasAvailable
    Progressing True    NewReplicaSetAvailable
  OldReplicaSets: <none>
  NewReplicaSet:  mcommerce-74d49b66fc (2/2 replicas created)
Events:
  Type      Reason          Age      From            Message
  ----      -----          ----      ----            -----
  Normal   ScalingReplicaSet 9m1s    deployment-controller  Scaled up replica set mcommerce-74d49b66fc to 2
```

Figure 3.19: Checking Deployments

The screenshot shows the deployment status, indicating successful pod creation.

3.5.4 Check ReplicaSet:

```
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ sudo kubectl get replicases
NAME          DESIRED   CURRENT   READY   AGE
mcommerce-74d49b66fc  2         2         2      11m
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ sudo kubectl describe replicases mcommerce-74d49b66fc
Name:           mcommerce-74d49b66fc
Namespace:      default
Selector:       app=mcommerce,pod-template-hash=74d49b66fc
Labels:         app=mcommerce
                pod-template-hash=74d49b66fc
Annotations:    deployment.kubernetes.io/desired-replicas: 2
                deployment.kubernetes.io/max-replicas: 3
                deployment.kubernetes.io/revision: 1
Controlled By: Deployment/mcommerce
Replicas:      2 current / 2 desired
Pods Status:   2 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=mcommerce
           pod-template-hash=74d49b66fc
  Containers:
    mcommerce:
      Image:      projdevopsiccn2/mcommerce:latest
      Port:       9001/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
      Node-Selectors: <none>
      Tolerations:  <none>
Events:
  Type      Reason     Age   From           Message
  ----      ----     --   --   -----
  Normal   SuccessfulCreate  11m   replicaset-controller  Created pod: mcommerce-74d49b66fc-cbplx
  Normal   SuccessfulCreate  11m   replicaset-controller  Created pod: mcommerce-74d49b66fc-g6bmb
```

Figure 3.20: Checking ReplicaSet

The screenshot verifies the ReplicaSet managing the pods.

3.5.5 Check All Components:

```
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo/.kubernetes$ sudo kubectl get all -o wide -n=product-local-use
NAME          READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED-NODE   READINESS   GATES
pod/mcommerce-74d49b66fc-f87n7  1/1    Running   0        4m53s  10.244.0.6  minikube  <none>        <none>
pod/mcommerce-74d49b66fc-z8fc7  1/1    Running   0        4m53s  10.244.0.5  minikube  <none>        <none>

NAME          TYPE        CLUSTER-IP      EXTERNAL-IP    PORT(S)    AGE   SELECTOR
service/product-service  ClusterIP  10.101.21.168  <none>        80/TCP    13m   app=mcommerce

NAME          READY   UP-TO-DATE   AVAILABLE   AGE   CONTAINERS   IMAGES          SELECTOR
deployment.apps/mcommerce  2/2     2           2          4m53s  mcommerce   projdevopsiccn2/mcommerce:latest  app=mcommerce

NAME          DESIRED   CURRENT   READY   AGE   CONTAINERS   IMAGES          SELECTOR
replicaset.apps/mcommerce-74d49b66fc  2         2         2      4m53s  mcommerce   projdevopsiccn2/mcommerce:latest  app=mcommerce,pod-template-hash=74d49b66fc
```

Figure 3.21: Displaying Resource Details

This screenshot displays detailed information about the resources in the Kubernetes cluster, such as deployments, pods, ReplicaSets, and services. It helps verify that all components are running correctly.

and provides useful information like pod IPs and node assignments.

```
zenlnex@vps-4f4e3327:/mcommerce/cicd-demo$ curl http://localhost:8001/api/v1/namespaces/product-local-use/services/product-service:80/proxy/ | jq
% Total    % Received % Xferd  Average Speed   Time     Time      Current
                                         Dload  Upload Total   Spent    Left  Speed
100  388    0  388    0     0  3629      0 --:--:-- --:--:-- --:--:-- 3695
[
  {
    "id": 1,
    "name": "Morocco Football Home Kit",
    "description": "Morocco 2022 World Cup Home Kit with Red and Green Colors",
    "price": 200
  },
  {
    "id": 2,
    "name": "Morocco Football Away Kit",
    "description": "Morocco 2022 World Cup Away Kit with White Color",
    "price": 200
  },
  {
    "id": 3,
    "name": "Morocco Football Pullover",
    "description": "Morocco 2022 World Cup Pullover with Red and Green Colors",
    "price": 500
  }
]
```

Figure 3.22: Accessing Data Hosted by the Product Microservice

This screenshot demonstrates that the data within the `mcommerce` microservice is accessible. It shows the successful connection to the service and retrieval of the data, confirming that the microservice is functioning as expected.

Chapter 4

Extension of the CI pipeline with GitOps

4.0.1 What is GitOps

Infrastructure as code give us the ability to describe our deployment environment using machine readable code. This help us to automate and manage infrastructure efficiently. GitOps leverages version control systems, like Git, to enable versioned and auditable workflows for managing code that's used for infrastructure and application deployments.

4.0.2 Using ArgoCD

Argo CD is a declarative, GitOps continuous delivery tool for Kubernetes. It enables users to manage and automate the deployment of applications in Kubernetes clusters by leveraging Git repositories as the source of truth for application states.

4.0.3 Preparing the IaC repository

So the first thing we did is creating a repository to hold our IaC.

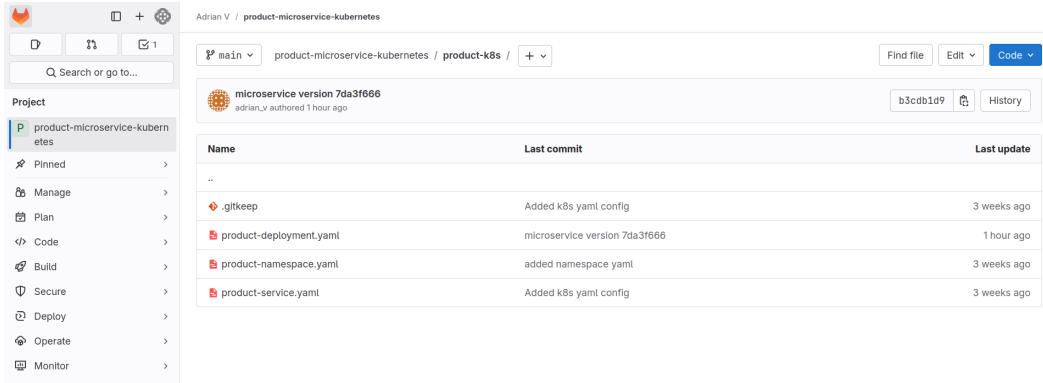


Figure 4.1

As we can see these are the same yaml file used in previous chapter.

The next important thing is to add a new job in our project CI pipeline to update the IaC repository to deploy the newly pushed docker image pushed by the CI pipeline.

```

48 gitops-k8s-deploy:
49   rules:
50     - if: $CI_COMMIT_BRANCH == "master"
51       changes:
52         paths:
53           - Dockerfile
54           - src/**/*
55       image: bitnami/git:latest
56       stage: deploy-k8s
57       before_script:
58         - git config --global user.email "62484408+terriblemoment0836@users.noreply.github.com"
59         - git config --global user.name "adrian_v"
60       script:
61         - git clone --depth 1 "https://gitlab-ci:$K8S_REPO_TOKEN@gitlab.com/adrian_v/product-microservice-kubernetes.git"
62         - cd product-microservice-kubernetes/product-k8s
63         - sed -i "s/projdevopsicn2.*$/projdevopsicn2/mcommerce:$TAG/g" product-deployment.yaml
64         - cat product-deployment.yaml
65         - git add product-deployment.yaml
66         - git commit -m "microservice version $TAG"
67         - git push -uf origin main
68

```

Figure 4.2

4.0.4 Argo CD configuration

Now we need to install Argo CD in our kubernetes cluster and point it to the IaC repository.

```

077R@vps-4f4e3327:~$ kubectl create namespace argocd
namespace/argocd created
077R@vps-4f4e3327:~$ kubectl apply -n argocd -f https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml
customresourcedefinition.apiextensions.k8s.io/applicationsets.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/applicationsets.argoproj.io created
customresourcedefinition.apiextensions.k8s.io/applications.argoproj.io created
serviceaccount/argocd-application-controller created
serviceaccount/argocd-applicationset-controller created
serviceaccount/argocd-dex-server created
serviceaccount/argocd-notifications-controller created
serviceaccount/argocd-redis created
serviceaccount/argocd-repo-server created
serviceaccount/argocd-server created
role.rbac.authorization.k8s.io/argocd-application-controller created
role.rbac.authorization.k8s.io/argocd-applicationset-controller created
role.rbac.authorization.k8s.io/argocd-dex-server created
role.rbac.authorization.k8s.io/argocd-notifications-controller created

```

Figure 4.3: Installing Argo CD

To access Argo CD web UI we use the command 'kubectl port-forward'

```
D77R@vps-4f4e3327:~$ kubectl port-forward svc/argocd-server -n argocd 8080:443
Forwarding from 127.0.0.1:8080 -> 8080
Forwarding from [::1]:8080 -> 8080
```

Figure 4.4: Port forwarding to access Argo CD

Then using ssh port forwarding we can access the UI locally.

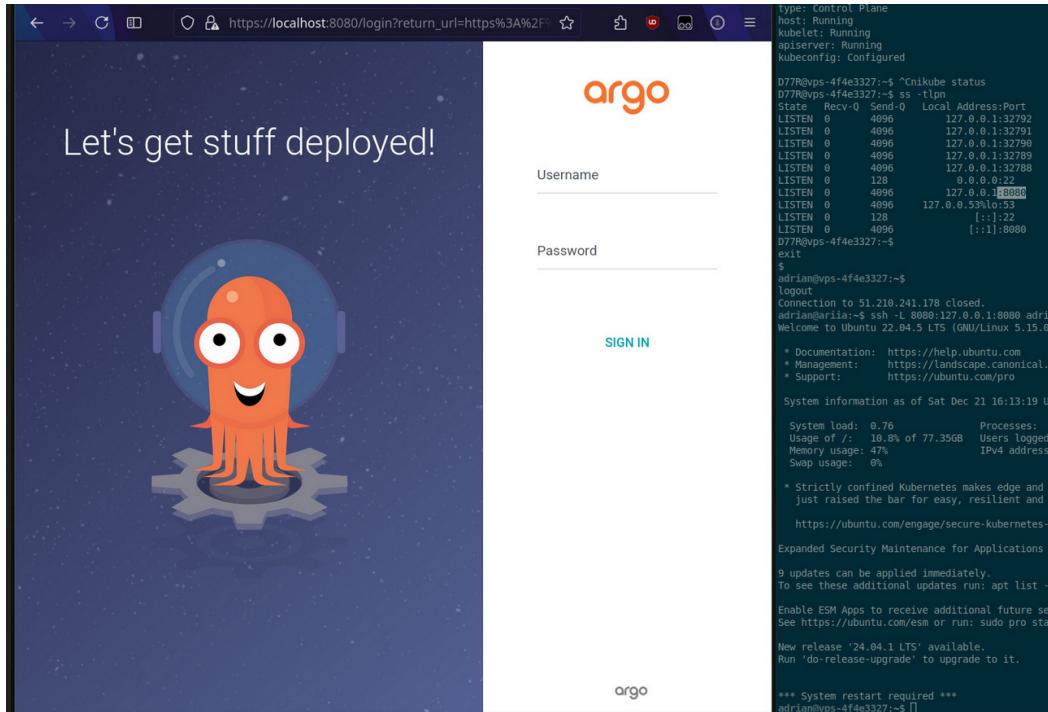


Figure 4.5: Ssh tunneling to access the UI locally

Then we add the IaC repository to Argo CD

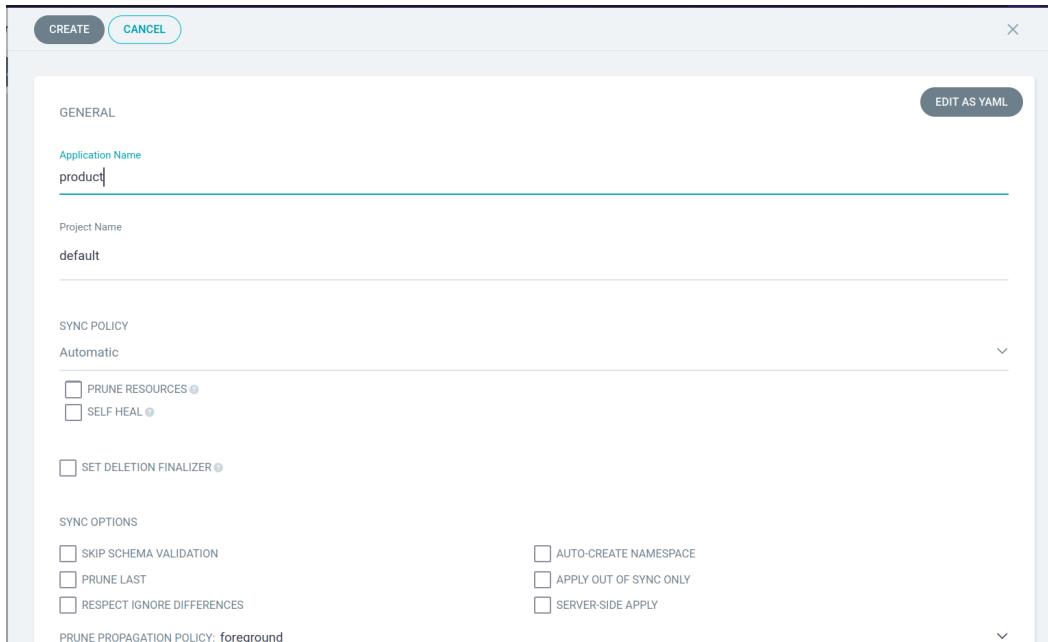


Figure 4.6

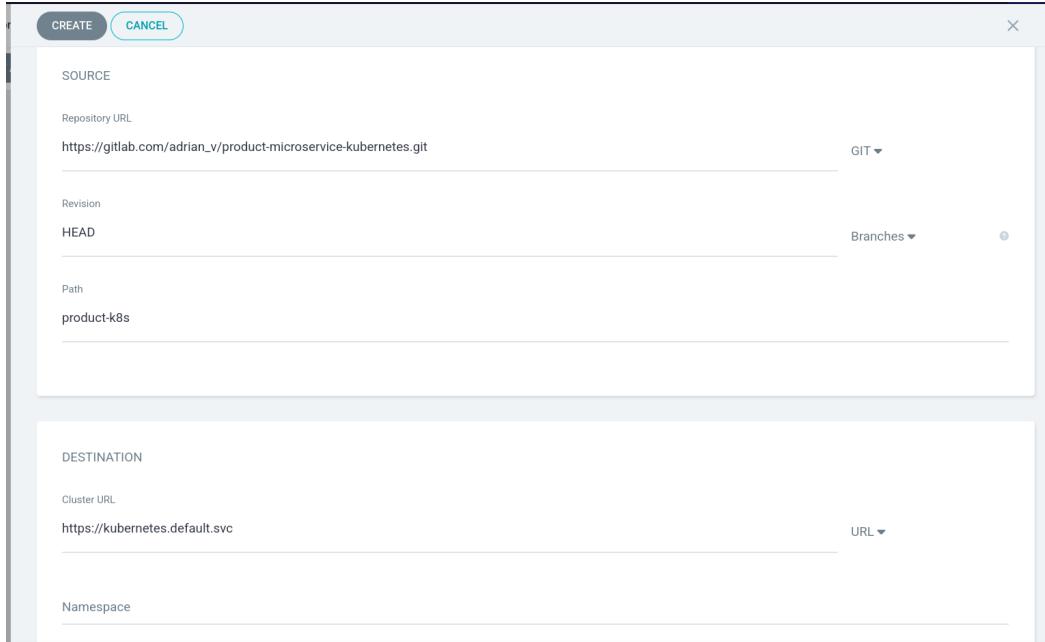


Figure 4.7

And as a result, we can see that our repository got synced, and the deployment is done.

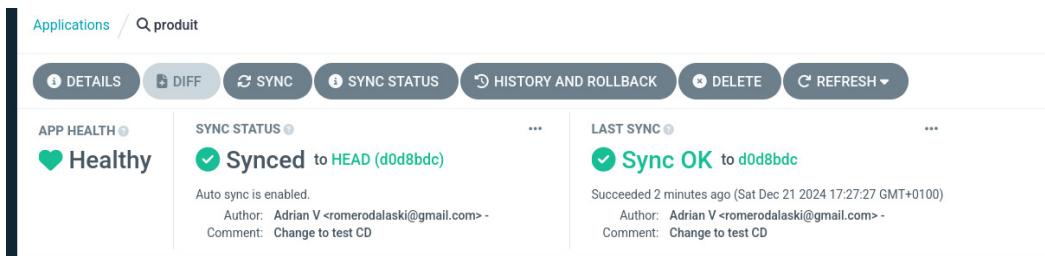


Figure 4.8



Figure 4.9

Chapter 5

Optimisation et Sécurisation du VCS et du CI Pipeline

5.1 Optimizing the CI pipeline

In this part we will implement some practices to help optimize our CI pipeline execution in terms of: Speed of execution, downloaded Data and configuration organization.

The main practices we will discuss are:

- Running CI inside a local GitLab Runner.
- Job rules.
- Caching.
- Generic jobs.

Next, we will discuss each practice, its benefits, and implementation.

5.1.1 CI in Local GitLab Runner

A GitLab Runner is an open-source application that is used to run jobs in a GitLab CI/CD (Continuous Integration/Continuous Deployment) pipeline. It executes scripts defined in a `.gitlab-ci.yml` file, allowing for automated testing, building, and deployment of code. By default GitLab provides its own GitLab runners that run jobs inside a docker container. One of the drawbacks of using the GitLab provided ones is:

- Resource Limitations:** GitLab's shared runners may have limited resources (CPU, memory, etc.), leading to slower job execution times, especially during peak usage periods.
- Queue Times:** Shared runners can experience longer queue times due to high demand, delaying the execution of CI/CD pipelines and impacting development cycles.
- Customization:** Running your own GitLab Runner allows for greater customization of the environment, including specific software versions, dependencies, and configurations required for your projects.
- Security and Compliance:** Using shared runners may pose security risks, as jobs run in a shared environment. Custom runners ensure sensitive data and proprietary code are handled securely.

Thus it's beneficial for us to run the CI pipeline in our own GitLab runner. Now to achieve that we will run a GitLab Runner in a Docker container (docker executor). So the first thing is to download the gitlab-runner image:

```
adrian@vps-4f4e3327:~$ sudo docker pull gitlab/gitlab-runner
Using default tag: latest
latest: Pulling from gitlab/gitlab-runner
Digest: sha256:11e3a44e7df21a32691a6975c7824c682787f742776d705ce89a74a780e3809f
Status: Image is up to date for gitlab/gitlab-runner:latest
docker.io/gitlab/gitlab-runner:latest
adrian@vps-4f4e3327:~$ sudo docker images | grep gitlab-runner
registry.gitlab.com/gitlab-org/gitlab-runner/gitlab-runner-helper    x86_64-v17.7.0      8e4bed5b3562   2 weeks ago   89.9MB
gitlab/gitlab-runner          latest                      6da9e8e16413   3 weeks ago   844MB
adrian@vps-4f4e3327:~$
```

Figure 5.1

Then we create a volume to persist our runner configuration.

```
adrian@vps-4f4e3327:~$ sudo docker volume create gitlab-runner-config
gitlab-runner-config
adrian@vps-4f4e3327:~$ sudo docker volume ls | grep runner-config
local     gitlab-runner-config
adrian@vps-4f4e3327:~$
```

Figure 5.2

Now we need to configure the runner by editing the file 'config.toml' inside the volume we created.

```
root@vps-4f4e3327:~# vim /var/lib/docker/volumes/gitlab-runner-config/_data/config.toml
```

Figure 5.3

```
concurrent = 4
check_interval = 0
shutdown_timeout = 0

[session_server]
# listen_address = "[::]:8093"
    session_timeout = 1800

[[runners]]
    name = "main-runner"
    url = "https://gitlab.com"
    executor = "docker"
    token = "xxxxxxxxxxxxxxxxxxxxxx"
[runners.docker]
    image = "docker:dind"
    privileged = true
    disable_cache = false
    cache_dir = ""
    volumes = ["/certs/client", "/cache"]
~
~
```

Figure 5.4

The provided configuration file sets up a GitLab Runner with the following key features:

- **Concurrency:** The runner can execute up to 4 jobs simultaneously.
- **Shutdown Behavior:** There is no timeout for shutting down the runner (shutdown timeout set to 0).
- **Session Management:** A session server is configured with a 30-minute session timeout.
- **Runner Details:**
 - The runner is named "main-runner."
 - It connects to the GitLab instance at <https://gitlab.com>.
 - It uses the Docker executor to run jobs.
 - A registration token is specified.
- **Docker Configuration:**
 - The runner uses the `docker:dind` image for Docker-in-Docker functionality.
 - It runs in privileged mode to allow Docker commands within the container.
 - Caching is enabled to speed up builds, and specific volumes are mounted for sharing data between jobs (artifacts).

Now to obtain the GitLab runner token we need to create a new GitLab runner in our GitLab Project.

Runners are processes that pick up and execute CI/CD jobs for GitLab. [What is GitLab Runner?](#)

Register as many runners as you want. You can register runners as separate users, on separate servers, and on your local machine.

How do runners pick up jobs?

Runners are either:

- **active** - Available to run jobs.
- **paused** - Not available to run jobs.

Tags control which type of jobs a runner can handle. By tagging a runner, you make sure runners only handle the jobs they are equipped to run. [Learn more](#).

Project runners

These runners are assigned to this project.

New project runner :

Instance runners

These runners are available to all groups and projects.

Each CI/CD job runs on a separate, isolated virtual machine.

Enable instance runners for this project

Figure 5.5

First we disable Giltab provided instance runners and we create a new project runner

Tags

Tags

Add tags to specify jobs that the runner can run. [Learn more](#).

Separate multiple tags with a comma. For example, `macos`, `shared`.

Run untagged jobs

Use the runner for jobs without tags in addition to tagged jobs.

Configuration (optional)

Runner description

Paused

Stop the runner from accepting new jobs.

Protected

Use the runner on pipelines for protected branches only.

Lock to current projects

Use the runner for the currently assigned projects only. Only administrators can change the assigned projects.

Maximum job timeout

Maximum amount of time the runner can run before it terminates. If a project has a shorter job timeout period, the job timeout period of the instance runner is used instead.

Enter the job timeout in seconds. Must be a minimum of 600 seconds.

Create runner

Figure 5.6

After we click 'Create runner' a token appears that we will need to put in the config.toml file.

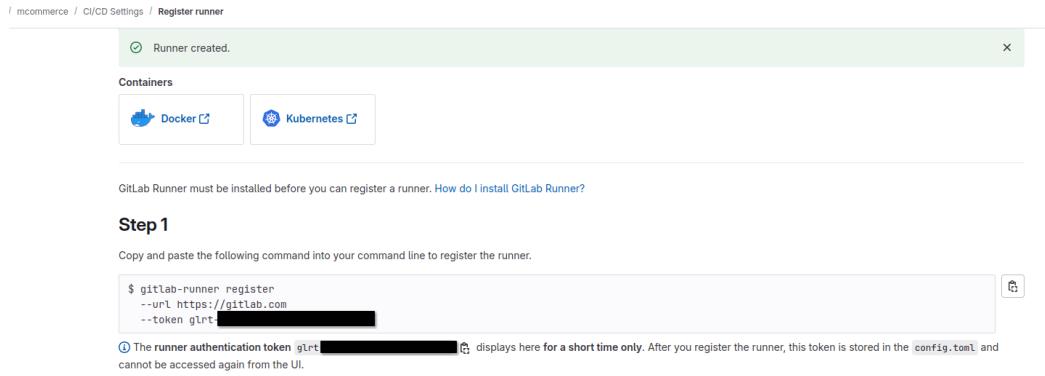


Figure 5.7

After this, we start the gitlab runner container and as we can see the runner is working.

```
root@vps-4f4e3327:~# docker ps | tail -1
4b8698da1c69  gitlab/gitlab-runner:latest          "/usr/bin/dumb-init ..."  3 weeks ago   Up 2 weeks
```

Figure 5.8: gitlab-runner container is running

```
root@vps-4f4e3327:~# docker inspect gitlab-runner | grep -i state -A 2
  "State": {
    "Status": "running",
    "Running": true,
```

Figure 5.9

Assigned project runners



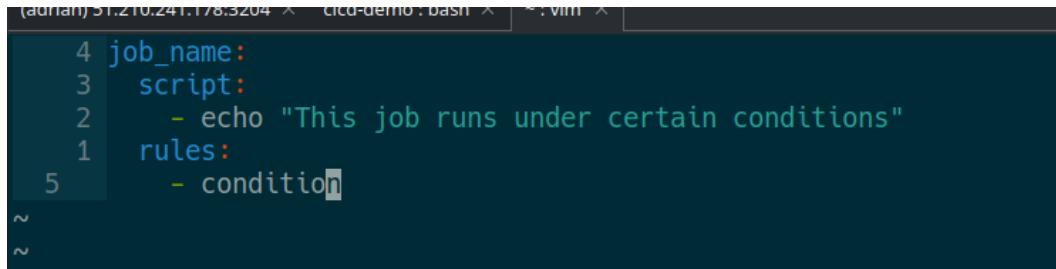
Figure 5.10: The local gitlab-runner is connected to GitLab

5.1.2 Job rules

Job rules are utilized to control the execution of jobs based on predefined conditions. These rules allow for the specification of criteria that determine when a job should be executed or skipped, thereby facilitating a more efficient workflow.

By leveraging job rules, the behavior of CI/CD processes can be tailored to respond dynamically to various contextual factors, such as the specific branch being built or the presence of certain environment variables. This capability enhances the adaptability of the pipeline and optimizes resource utilization by ensuring that jobs are executed only when relevant.

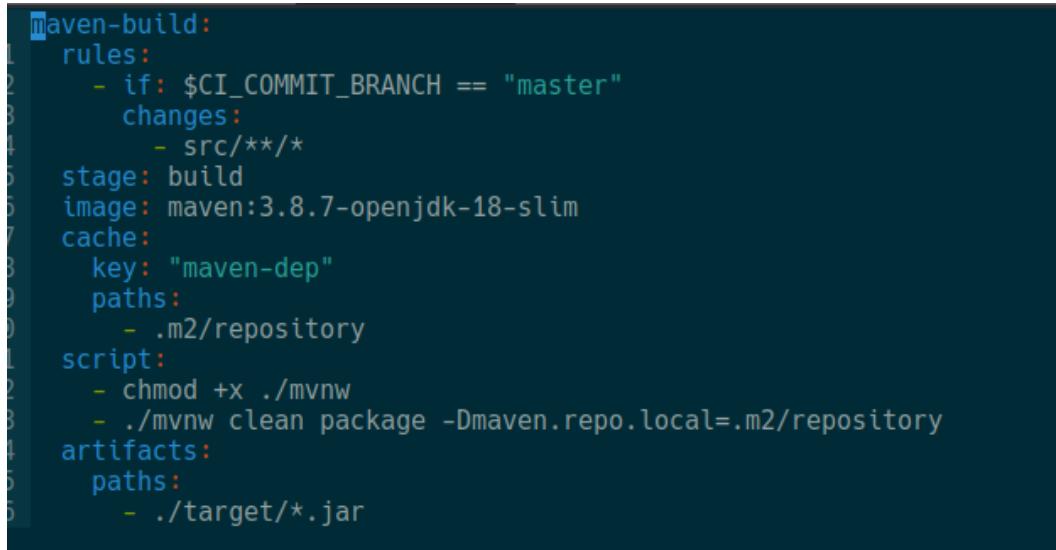
In GitLab CI/CD, YAML job rules are defined using the 'rules' keyword within a job's configuration in the .gitlab-ci.yml file. The syntax allows for specifying conditions under which a job should run or be skipped. Here is a simple example:



```
4 job_name:
3   script:
2     - echo "This job runs under certain conditions"
1   rules:
5     - condition
```

Figure 5.11

In our case, we defined several rules for several jobs across the pipeline:



```
maven-build:
  rules:
    - if: $CI_COMMIT_BRANCH == "master"
      changes:
        - src/**/*
  stage: build
  image: maven:3.8.7-openjdk-18-slim
  cache:
    key: "maven-dep"
    paths:
      - .m2/repository
  script:
    - chmod +x ./mvnw
    - ./mvnw clean package -Dmaven.repo.local=.m2/repository
  artifacts:
    paths:
      - ./target/*.jar
```

Figure 5.12

```

1 .docker-build:
2   rules:
3     - if: $CI_COMMIT_BRANCH == "master"
4       changes:
5         paths:
6           - Dockerfile
7           - src/**/*
8   stage: container
9   image: docker
10  dependencies:
11    - maven-build
12  services:
13    - docker:dind
14  before_script:
15    - docker login -u $DOCKER_LOGIN -p $DOCKER_TOKEN
16  script:
17    - docker version
18    - docker build -t $DOCKER_LOGIN/$IMAGE_NAME:$TAG .
19    - docker push $DOCKER_LOGIN/$IMAGE_NAME:$TAG

```

Figure 5.13

```

2 .gitops-k8s-deploy:
1   variables:
2     GIT_EMAIL: ""
3     GIT_USERNAME: ""
4     GITLAB_TOKEN_NAME: ""
5     GITLAB_TOKEN: ""
6     GITLAB_REPO_URL: ""
7     MICROSERVICE_NAME: ""
8
9   rules:
10    - if: $CI_COMMIT_BRANCH == "master"
11      changes:
12        paths:
13          - Dockerfile
14          - src/**/*
15    image: bitnami/git:latest
16    stage: deploy-k8s
17    before_script:
18      - git config --global user.email $GIT_EMAIL
19      - git config --global user.name $GIT_USERNAME
20      - git clone --depth 1 "https://$GITLAB_TOKEN_NAME:$GITLAB_TOKEN@$GITLAB_REPO_URL"
21    script:
22      - cd $MICROSERVICE_NAME-microservice-kubernetes/$MICROSERVICE_NAME-k8s
23      - sed -i "s/projdevopsicn2.*$/projdevopsicn2\`mcommerce:$TAG/g" $MICROSERVICE_NAME-deployment.yaml
24      - cat $MICROSERVICE_NAME-deployment.yaml
25      - git add $MICROSERVICE_NAME-deployment.yaml
26      - git commit -m "microservice version $TAG"
27      - git push -uf origin main

```

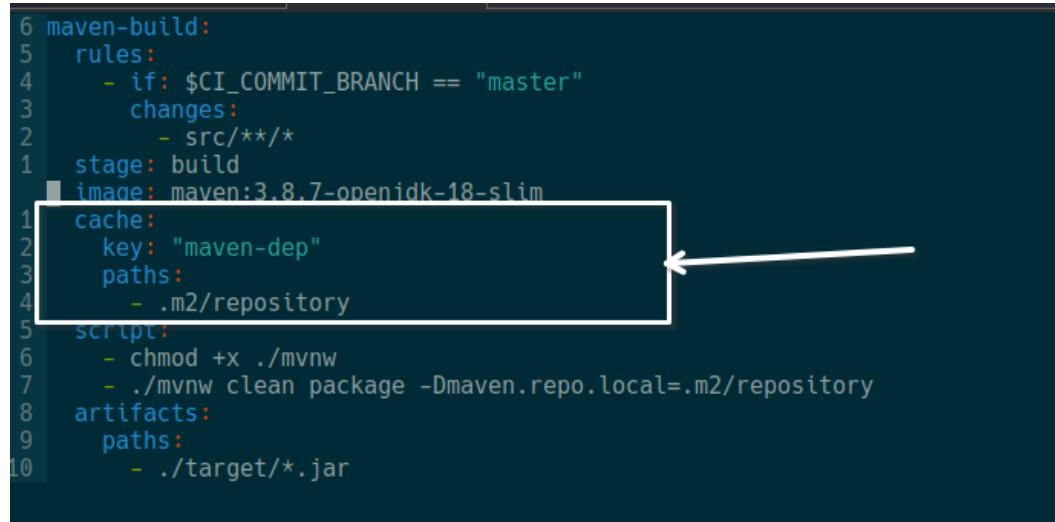
Figure 5.14

- In figure 5.12:** The rule here signifies that we want to build the project only if changes happened to the master branch and these changes affect only the application source code inside the src directory. thus, we avoid unnecessary builds when other files than changes.
- in figure 5.13:** The rule here signifies that we want to containerize the application only if changes happened to the 'Dockerfile' or the application source code in the master branch.
- In figure 5.14:** The rule here signifies that we want to update our IaC repository only if there are changes to the way the application is containerized or there are changes to the source code of the application in the master branch.

5.1.3 Caching

Caching is defined as a mechanism by which frequently accessed data is stored temporarily to improve the performance of subsequent operations. By utilizing caching, the retrieval of data is expedited, thereby reducing the time and resources required for tasks such as job execution in CI/CD pipelines. Caching is implemented to minimize redundant data processing and to enhance overall efficiency within the development workflow.

In our case, caching makes sense only for the application dependencies. Thus, we modified the maven-build job to cache downloaded dependencies to be used for subsequent builds.



```

6 maven-build:
5   rules:
4     - if: $CI_COMMIT_BRANCH == "master"
3       changes:
2         - src/**/*
1       stage: build
      image: maven:3.8.7-openjdk-18-slim
1       cache:
2         key: "maven-dep"
3         paths:
4           - .m2/repository
5       script:
6         - chmod +x ./mvnw
7         - ./mvnw clean package -Dmaven.repo.local=.m2/repository
8       artifacts:
9         paths:
10        - ./target/*.jar

```

Figure 5.15

The modification includes adding a new entry 'cache' that declares that we want to cache the content of the directory '.m2/repository' that contains the downloaded dependencies. Also, we gave this cache a key name as a way to identify it. Note that we have changed the maven command (-Dmaven.repo.local) to specify to it to use the '.m2/repository' because it's the cached directory.

5.1.4 Generic jobs

In GitLab, the generic job include feature is defined as a mechanism that allows the inclusion of external YAML files into a CI/CD pipeline configuration. By utilizing this feature, shared configurations can be centralized, promoting consistency across multiple projects. The benefits of using job includes include improved maintainability, as updates can be made in one location and automatically reflected in all pipelines, and enhanced collaboration, as teams can share common job definitions easily. Additionally, the overall complexity of pipeline configurations is reduced, leading to clearer and more manageable

CI/CD setups.

In our case, we created a new directory '.gitlab' each file represents a generic job, That we will include in the main 'gitlab-ci.yml' file and apply changes specific to the current project, as shown in figure 5.18.

The screenshot shows a GitLab repository named 'cicd-demo'. In the '.gitlab' directory, there are several files listed:

Name	Last commit	Last update
...		
docker-build.yml	more templating	2 weeks ago
k8s-deploy.yml	more templating	2 weeks ago
maven-build.yml	splitting .gitlab-ci.yml	2 weeks ago
stages.yml	splitting .gitlab-ci.yml	2 weeks ago
variables.yml	splitting .gitlab-ci.yml	2 weeks ago

Figure 5.16: Different generic jobs

```

/docker-build:
  variables:
    DOCKER_LOGIN: "user"
    IMAGE_NAME: "image"
  rules:
    - if: $CI_COMMIT_BRANCH == "master"
      changes:
        paths:
          - Dockerfile
          - src/**/*
  stage: container
  image: docker
  services:
    - docker:dind
  before_script:
    - docker login -u $DOCKER_LOGIN -p $DOCKER_TOKEN
  script:
    - docker version
    - docker build -t $DOCKER_LOGIN/$IMAGE_NAME:$TAG .
    - docker push $DOCKER_LOGIN/$IMAGE_NAME:$TAG

```

Figure 5.17: Example of a generic job with default variable that will be overridden

```

include:
  - '.gitlab/stages.yml'
  - '.gitlab/variables.yml'
  - '.gitlab/maven-build.yml'
  - '.gitlab/docker-build.yml'
  - '.gitlab/k8s-deploy.yml'

docker-build:
  variables:
    DOCKER_LOGIN: "projdevopsiccn2"
    IMAGE_NAME: "mcommerce"
  extends: .docker-build

gitops-k8s-deploy:
  variables:
    GIT_EMAIL: "62484408+terriblemoment0836x@users.noreply.github.com"
    GIT_USERNAME: "adrian_v"
    GITLAB_TOKEN_NAME: "gitlab-ci"
    GITLAB_TOKEN: "$K8S_REPO_TOKEN"
    GITLAB_REPO_URL: "gitlab.com/adrian_v/product-microservice-kubernetes.git"
    MICROSERVICE_NAME: "product"
  extends: .gitops-k8s-deploy

```

Figure 5.18: Main CI configuration where the included generic jobs are customized

5.2 gestion des groupes et permissions

5.2.1 Définir des rôles

The role definition section establishes two primary access levels within the project: maintainer and developer roles.

The maintainer role is granted elevated privileges necessary for project administration and security management, including code deployment, environment configuration, and security policy enforcement.

The developer role is configured with permissions appropriate for daily development tasks while maintaining security boundaries.

```

1  # Role definitions
2  roles:
3    maintainer_role:
4      permissions:
5        - push_code
6        - merge_code
7        - manage_project_settings
8        - manage_members
9        - create_protected_tags
10       - delete_protected_tags
11       - manage_pipelines
12       - manage_environments
13       - manage_deployments
14       - approve_merge_requests
15       - override_project_settings
16       - manage_security_policies
17
18   developer_role:
19     permissions:
20       - push_code
21       - create_merge_requests
22       - create_branches
23       - run_pipelines
24       - create_issues
25       - manage_issues
26       - approve_merge_requests
27       - deploy_to_staging
28

```

Figure 5.19: Defining Roles

```

# Role definitions
roles:
  maintainer_role:
    permissions:
      - push_code
      - merge_code
      - manage_project_settings

```

- manage_members
- create_protected_tags
- delete_protected_tags
- manage_pipelines
- manage_environments
- manage_deployments
- approve_merge_requests
- override_project_settings
- manage_security_policies

developer_role :

permissions :

- push_code
- create_merge_requests
- create_branches
- run_pipelines
- create_issues
- manage_issues
- approve_merge_requests
- deploy_to_staging

5.2.2 Group definitions with roles

Group definitions implement the practical application of roles through structured team organization.

This configuration creates two distinct groups: project-maintainers and project-developers. Each group is assigned specific visibility settings and access levels, ensuring proper segregation of duties and access control. The private visibility setting ensures that group membership and activities remain restricted to authorized personnel.

```

46 # Group definitions
29 groups:
30   - name: project-maintainers
31     path: project-maintainers
32     description: "Equipe de maintenance avec accès complet au projet"
33     visibility: private
34     access_level: maintainer
35
36   - name: project-developers
37     path: project-developers
38     description: "Equipe de développement"
39     visibility: private
40     access_level: developer
41
42

```

Figure 5.20: Groups

groups :

- name: project-maintainers
 - path: project-maintainers
 - description: "Equipe de maintenance avec accès complet au projet"
 - visibility: private
 - access_level: maintainer
 - role: maintainer_role
- members :
 - username: maintainer1
 - access_level: maintainer

- name: project-developers
 - path: project-developers
 - description: "Equipe de développement"
 - visibility: private
 - access_level: developer
 - role: developer_role
- members :
 - username: developer1
 - access_level: developer

5.2.3 Project access configuration

5.2.4 Protection des branches

Branch protection mechanisms safeguard critical code bases by implementing hierarchical access controls. The main branch receives the highest level of protection, requiring maintainer approval for all modifications.

The develop branch allows developer access while maintaining maintainer oversight for critical operations.

```

42
43 # Project access configuration
44 project:
45   # Protection des branches
46   protected_branches:
47     - name: main
48       push_access_level: maintainer
49       merge_access_level: maintainer
50       unprotect_access_level: maintainer
51
52     - name: develop
53       push_access_level: developer
54       merge_access_level: developer
55       unprotect_access_level: maintainer
56

```

Figure 5.21: protection des branches

project :

```

# Protection des branches

protected_branches:

  - name: main
    push_access_level: maintainer
    merge_access_level: maintainer
    unprotect_access_level: maintainer

  - name: develop
    push_access_level: developer
    merge_access_level: developer
    unprotect_access_level: maintainer

```

5.2.5 Configuration des merge requests

The merge request configuration implements a code review process that ensures quality control while maintaining development efficiency. A single approval is required from either maintainers or developers, promoting collaborative code review while preventing deployment bottlenecks.

```

44 project:
45   # Configuration des merge requests
46   merge_requests:
47     approval_rules:
48       - name: "Default"
49         approvals_required: 1
50         user_groups:
51           - project-maintainers
52           - project-developers
53
54
55

```

Figure 5.22: merge request

```
# Configuration des merge requests

merge_requests:

  approval_rules:
    - name: "Default"
      approvals_required: 1

  user_groups:
    - project-maintainers
    - project-developers
```

5.2.6 CI/CD

The CI/CD configuration establishes a secure deployment pipeline by implementing environment-specific access controls. Production deployments are restricted to maintainers only, ensuring maximum security for critical infrastructure. Staging environment access is granted to developers, facilitating testing and validation while maintaining the security boundary between development and production environments.

```
44 project:
45
46   # Paramètres CI/CD
47   ci_cd:
48     protected_environments:
49       - name: production
50         deploy_access_levels:
51           - group_id: project-maintainers
52             access_level: maintainer
53
54       - name: staging
55         deploy_access_levels:
56           - group_id: project-developers
57             access_level: developer
58
```

Figure 5.23: CI/CD parameter

Paramètres CI/CD

```
ci_cd:
  protected_environments:
    - name: production
      deploy_access_levels:
        - group_id: project-maintainers
          access_level: maintainer

    - name: staging
      deploy_access_levels:
        - group_id: project-developers
          access_level: developer
```

```

deploy_access_levels:
  - group_id: project-developers
    access_level: developer

```

5.2.7 Configuration des membres

Member configuration implements the practical assignment of users to their respective groups and access levels. This section defines specific user-role associations, ensuring that each team member has appropriate access permissions for their responsibilities.

```

44 project:
78
79 # Configuration des membres
80 members:
81   - group: project-maintainers
82     users:
83       - username: soufian-elouazzani1
84         access_level: maintainer
85
86   - group: project-developers
87     users:
88       - username: adrian_v
89         access_level: developers

```

Figure 5.24: members

```

# Configuration des membres

members:
  - group: project-maintainers
    users:
      - username: soufian-elouazzani1
        access_level: maintainer

  - group: project-developers
    users:
      - username: adrian_v
        access_level: developer

finally

```



Figure 5.25: push the file to the root folder of the project

5.3 Protected Branches

In GitLab, permissions are fundamentally defined around the idea of having read or write permission to the repository and branches. To impose further restrictions on certain branches, they can be protected.

5.3.1 Protected Branch Controls

A protected branch controls:

- Which users can merge into the branch.
- Which users can push to the branch.
- If users can force push to the branch.

5.3.2 Who Can Modify a Protected Branch

When a branch is protected, the default behavior enforces these restrictions on the branch.

Action	Who can do it
Protect a branch	At least the Maintainer role.
Push to the branch	Anyone with Allowed permission. (1)
Force push to the branch	No one.
Delete the branch	No one. (2)

5.3.3 Steps to Protect a Branch for All Projects in a Group

- a) On the left sidebar, select **Search** or go to and find your group.
- b) Select **Settings > Repository**.
- c) Expand **Protected branches**.
- d) Select **Add protected branch**.
- e) From the **Allowed to merge** list, select a role that can merge into this branch.
- f) From the **Allowed to push and merge** list, select a role that can push to this branch.
- g) Select **Protect**.

Branch rule details

Rule target

master
1 matching branch

[Edit](#) [Delete rule](#)

Protect branch

Keep stable branches secure and force developers to use merge requests. [What are protected branches?](#)

Allowed to merge 1

Roles Maintainers

[Edit](#)

Allowed to push and merge 1

Changes require a merge request. The following users can push and merge directly.

Roles Maintainers

[Edit](#)

Allow force push



Require code owner approval



The tag `protected` is put on all protected branches.

Branch rules

Define rules for who can push, merge, and the required approvals for each branch. [Leave feedback.](#)

Branch rules 1

Add branch rule ▾

master default protected

View details

- Allowed to merge: Maintainers
- Allowed to push and merge: Maintainers

5.4 Add a Coverage Check Approval Rule

5.4.1 Steps to Add a Coverage-Check Approval Rule

- Go to your project and select **Settings > Merge requests**.
- Under **Merge request approvals**, do one of the following:
 - Next to the **Coverage-Check approval rule**, select **Enable**.
 - For manual setup, select **Add approval rule**, then enter the **Rule name**. For example: *Coverage Check*.

Add approval rule X

Rule name

←

Examples: QA, Security.

Target branch

▼

Apply this approval rule to all branches or a specific protected branch.

Required number of approvals

^

Users 1

 **marouane**
@marouane.tanafatt Delete

Groups 0

Project groups ▼

No groups have been added.

c) Select a **Target branch**.

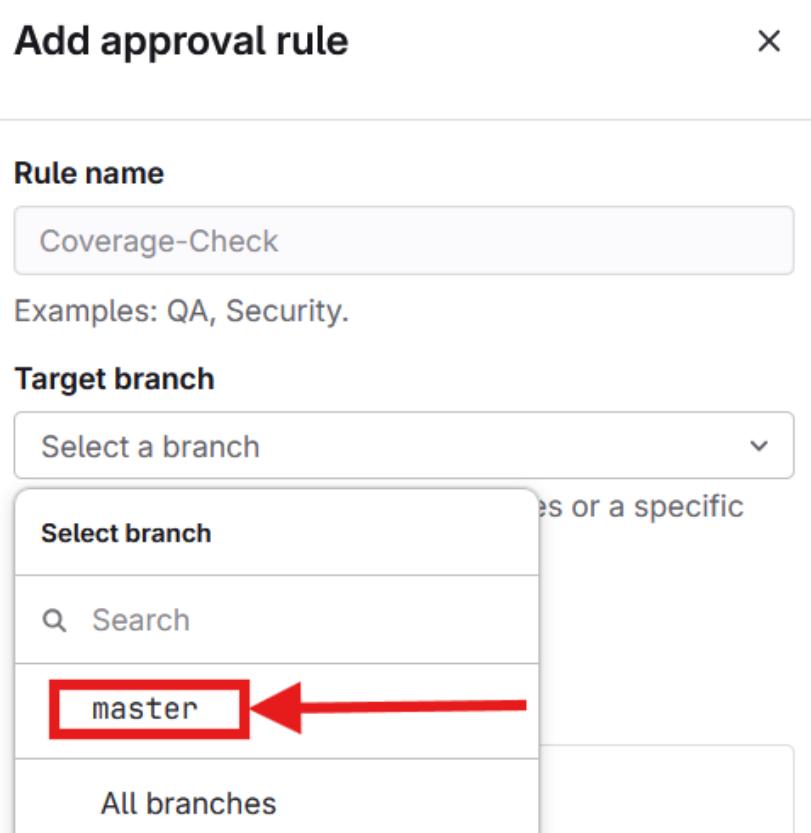
Add approval rule

Rule name
Coverage-Check

Examples: QA, Security.

Target branch
Select a branch

Select branch
Search
master
All branches



d) Set the **Required number of approvals**.

Add approval rule

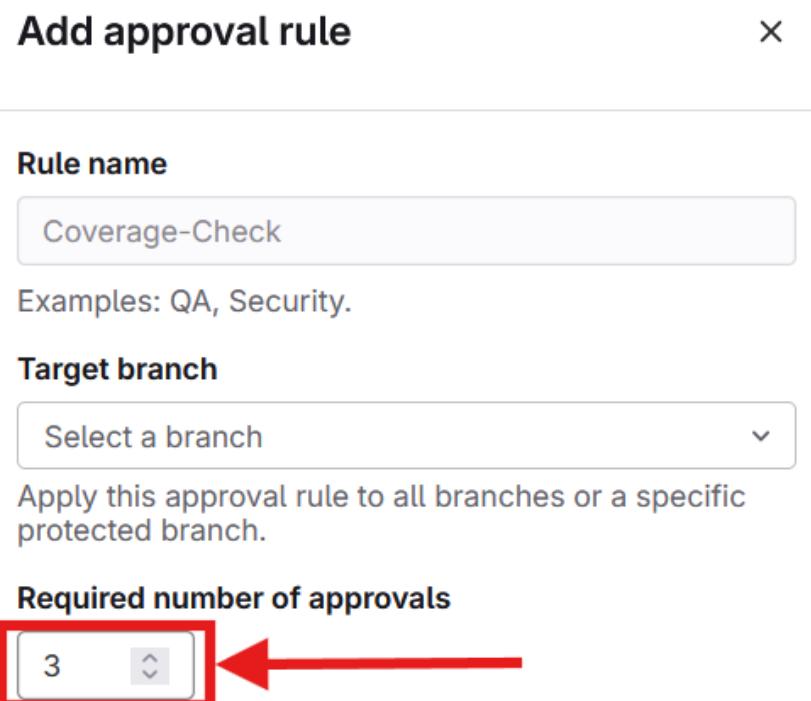
Rule name
Coverage-Check

Examples: QA, Security.

Target branch
Select a branch

Apply this approval rule to all branches or a specific protected branch.

Required number of approvals
3



e) Select the **Users or Groups** to provide approval.

Add approval rule X

Rule name

Coverage-Check

Examples: QA, Security.

Target branch

master ▼

Apply this approval rule to all branches or a specific protected branch.

Required number of approvals

3 ▼

Users 1

Search to add users

marouane
@marouane.tanafatt

Delete

Groups 0

Search to add group Project groups ▼

No groups have been added.

f) Select **Save changes**.

Merge request approvals

Define approval rules and settings to ensure separation of duties for new merge requests. [Learn more](#).

Approval rules 8 × 2		Add approval rule		
Approvers	Target branch	Approvals required	Actions	
Minimum required approvals	Any eligible user ?	All branches	0 Delete	
Coverage-Check ?		master	3 Delete	

5.5 Pre-Commit Hooks

5.5.1 Integration of Pre-Commit Hooks in GitLab

In this section, we address the issue raised in question 5-a):

« In the case where a developer disables pre-commit hooks locally, integrate a step in the pipeline to handle local commits without active pre-commit hooks. »

To address this, we set up two files:

- `.gitlab/pre-commit.yml` (or `pre-commit.yml`)
- `.pre-commit-config.yaml`

Main Objective: To enforce quality checks (such as Checkstyle, Spotless, dependency analysis, etc.) in the GitLab pipeline, even if a developer has disabled their local *pre-commit* hook.

5.5.2 Structure of the `pre-commit.yml` File in GitLab

The `pre-commit.yml` file (included later in `.gitlab-ci.yml`) defines a specific job that uses the `pre-commit` tool within the pipeline. Here is an example configuration snippet:

 pre-commit hooks and Gitlab CI integration
D77R authored 2 weeks ago

 [pre-commit.yml](#)  284 B

```

1 .pre-commit-check:
2   stage: test
3   image: python:3.9
4   cache:
5     paths:
6       - .cache/pre-commit
7   before_script:
8     - pip install pre-commit
9   script:
10    - pre-commit run --all-files || git add . || true
11   artifacts:
12     paths:
13       - pre-commit.log
14

```

Figure 5.26: Excerpt from the .gitlab/pre-commit.yml file

Explanation:

- **stage: test** : The job is associated with the testing phase of the pipeline.
- **image: python:3.9** : We use a Docker image containing Python to install and run `pre-commit`.
- **before_script** : Installation of `pre-commit` via `pip`.
- **script** : Execution of `pre-commit` hooks on all files (`-all-files`). *Even if a developer disables the hook locally, this job ensures the same checks are run in GitLab.*
- **artifacts** : The `pre-commit.log` file is retained for later review.

5.5.3 Configuration of `.pre-commit-config.yaml`

The `.pre-commit-config.yaml` file specifies the list of hooks to execute. Below is an example:

```

repos:
  - repo: https://github.com/pre-commit/pre-commit-hooks
    rev: v4.4.0
    hooks:
      - id: trailing-whitespace
      - id: end-of-file-fixer
      - id: check-yaml
      - id: detect-private-key #ICCCN o kda o hadi o mandiruch security 7chuma

  - repo: local
    hooks:
      - id: checkstyle
        name: Run Checkstyle
        entry: mvn checkstyle:check
        language: system
        always_run: true
      - id: spotless
        name: Format Java Code with Spotless
        entry: mvn spotless:apply
        language: system
        always_run: true
      - id: dependency-check
        name: Run OWASP Dependency Check
        entry: mvn org.owasp:dependency-check-maven:check
        language: system
        always_run: true
~
```

Figure 5.27: Excerpt from the .pre-commit-config.yaml file

Notes:

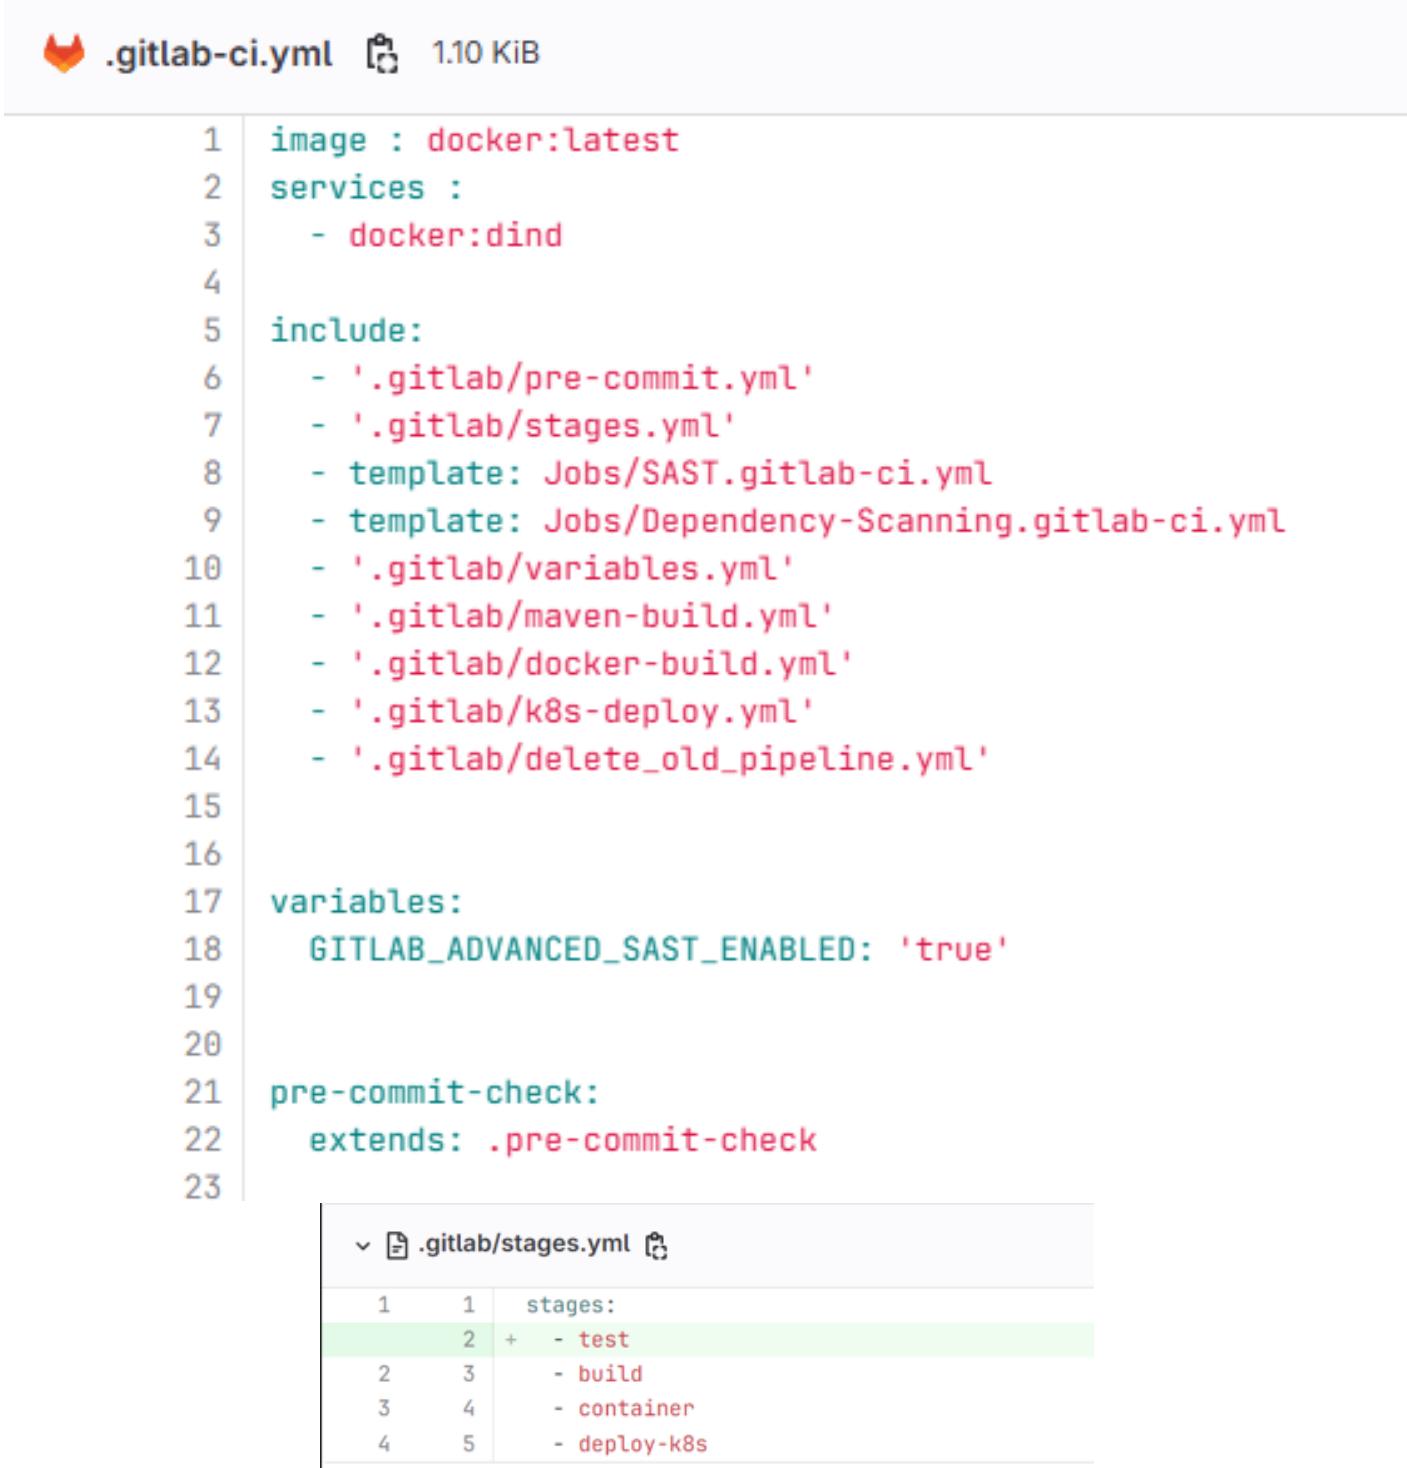
- **The official repository** (<https://github.com/pre-commit/pre-commit-hooks>), which provides common hooks such as:
 - `trailing-whitespace` (removal of unnecessary spaces),
 - `end-of-file-fixer` (ensuring newline at the end of files),
 - `check-yaml` (validation of YAML syntax),
 - `detect-private-key` (detection of potentially exposed private keys).
- **The local repository** (`repo: local`), which contains custom hooks based on Maven commands (`language: system`), including:
 - `checkstyle` (Java code analysis via `mvn checkstyle:check`),
 - `spotless` (automatic formatting with `mvn spotless:apply`),
 - `dependency-check` (vulnerability detection with `mvn org.owasp:dependency-check-maven:check`).

For each hook, the `always_run: true` option forces systematic execution on every commit, ensuring a consistent level of quality and security, even if the relevant files have not been modified. Overall, this

configuration guarantees code style compliance, adherence to security best practices, and continuous dependency control to prevent vulnerabilities.

5.5.4 Inclusion in `.gitlab-ci.yml` and Resolution of the Issue

To ensure this job runs *even if the developer has disabled their local pre-commit hook*, we include `pre-commit.yml` in `.gitlab-ci.yml`:



```

1  image : docker:latest
2  services :
3    - docker:dind
4
5  include:
6    - '.gitlab/pre-commit.yml'
7    - '.gitlab/stages.yml'
8    - template: Jobs/SAST.gitlab-ci.yml
9    - template: Jobs/Dependency-Scanning.gitlab-ci.yml
10   - '.gitlab/variables.yml'
11   - '.gitlab/maven-build.yml'
12   - '.gitlab/docker-build.yml'
13   - '.gitlab/k8s-deploy.yml'
14   - '.gitlab/delete_old_pipeline.yml'
15
16
17 variables:
18   GITLAB_ADVANCED_SAST_ENABLED: 'true'
19
20
21 pre-commit-check:
22   extends: .pre-commit-check
23

```

↳ **.gitlab/stages.yml**

1	1	stages:
2	2	+ - test
2	3	- build
3	4	- container
4	5	- deploy-k8s

Figure 5.28: Excerpt from the .gitlab-ci.yml file

- `include: '- '.gitlab/pre-commit.yml'` : Imports the dedicated configuration for `pre-commit` hooks, defined in a separate file.
- `pre-commit-check: extends: .pre-commit-check` : Creates a `pre-commit-check` job in the pipeline, inheriting the instructions contained in `.pre-commit-check`. This ensures the execution

of pre-commit checks on GitLab's side, even if a developer has disabled their hooks locally.

iccn / mcommerce / Pipelines / #1605234562

Testing post-commit hook

Passed D77R 3L1 created pipeline for commit 7de76595 7 minutes ago, finished 5 minutes ago

For master

latest ⚡ 4 jobs ⏲ 1 minute 58 seconds, queued for 1 seconds

Pipeline Jobs 4 Tests 0

Group jobs by Stage Job dependencies



Figure 5.29: GitLab pipeline validated after execution of `pre-commit-check`

Although developers can disable pre-commit hooks locally, the `pre-commit-check` job (figure 5.29, `test` phase) enforces the same checks on GitLab. Thus, even if the hook is disabled locally, the code must pass style and security checks for the pipeline to succeed. If issues are detected (trailing spaces, vulnerabilities, etc.), the pipeline fails and blocks merging or deployment. In the example above, all four jobs, including `pre-commit-check`, completed successfully, proving that the code passed these mandatory checks.

How Does This Solve the Problem?

- Developer disables the hook locally:** If, for any reason, the developer removes or disables `pre-commit` on their machine, they can push their code without the checks being performed locally.
- Mandatory execution in CI:** However, when creating a `merge request` or pushing, GitLab triggers the pipeline where the `pre-commit-check` job is defined.
- Validation or blocking:** If the code violates the rules defined in `.pre-commit-config.yaml`, the job fails, and the pipeline is blocked. The developer is then required to fix these issues (or justify any potential workaround) before merging their code.

Thus, the setup of the two files `pre-commit.yml` and `.pre-commit-config.yaml`, coupled with their execution in the GitLab pipeline (`.gitlab-ci.yml`), ensures compliance with the same quality rules, even if local hooks are disabled.

5.5.5 Evaluation of Using Frameworks like `pre-commit` (Tests for *pre-push*, *post-commit*)

After configuring `pre-commit` hooks, it is often relevant to extend `pre-commit` to include other stages of the Git lifecycle, such as *pre-push* and *post-commit*.

- ***pre-push***: Triggered before pushing code. Allows rerunning essential checks (unit tests, security analysis) to block non-compliant pushes.
- ***post-commit***: Executed immediately after creating a commit. Ideal for running internal scripts (e.g., updating documentation) without blocking the commit phase.

```
D77R@vps-4f4e3327:~/cicd-demo/.git/hooks$ ls
applypatch-msg.sample      pre-commit           pre-receive.sample
commit-msg.sample          pre-commit.sample    prepare-commit-msg.sample
fsmonitor-watchman.sample pre-merge-commit.sample push-to-checkout.sample
post-commit                pre-push             update.sample
post-update.sample         pre-push.sample
pre-applypatch.sample      pre-rebase.sample

D77R@vps-4f4e3327:~/cicd-demo$ pre-commit run --all-files
trim trailing whitespace.....Passed
fix end of files.....Passed
check yaml.....Passed
detect private key.....Passed
Run Checkstyle.....Passed
Format Java Code with Spotless.....Passed
Run OWASP Dependency Check.....Passed
```

Figure 5.30: `pre-commit run -all-files` executing various checks (formatting, checkstyle, etc.)

Benefits of `pre-commit` for *pre-push* or *post-commit*:

- **Standardize checks**: Reuse the same logic for checks (linting, tests, etc.) during push or immediately after commit.
- **Reduce risks of broken commits**: Prevent untested code from being pushed to the remote branch.

- **Improve feedback:** If *pre-push* hooks detect an issue, the developer is notified before polluting the repository.

```
D77R@vps-4f4e3327:~/cicd-demo$ vim .gitlab-ci.yml
D77R@vps-4f4e3327:~/cicd-demo$ vim .pre-commit-config.yaml
D77R@vps-4f4e3327:~/cicd-demo$ git add .pre-* .gitlab-ci.yaml
fatal: pathspec '.gitlab-ci.yaml' did not match any files
D77R@vps-4f4e3327:~/cicd-demo$ git add .pre-* .gitlab-ci.yml
D77R@vps-4f4e3327:~/cicd-demo$ git commit -m "pre-commit hooks and Gitlab CI integration"
[WARNING] Unstaged files detected.
[INFO] Stashing unstaged files to /home/D77R/.cache/pre-commit/patch1735407789-3517086.
trim trailing whitespace.....Passed
fix end of files.....Passed
check yaml.....Passed
black.....(no files to check)Skipped
[INFO] Restored changes from /home/D77R/.cache/pre-commit/patch1735407789-3517086.
[master eefff7d8] pre-commit hooks and Gitlab CI integration
 2 files changed, 28 insertions(+), 4 deletions(-)
   create mode 100644 .pre-commit-config.yaml
D77R@vps-4f4e3327:~/cicd-demo$ git push origin master
Username for 'https://gitlab.com': D73L1
Password for 'https://D73L1@gitlab.com':
Enumerating objects: 6, done.
Counting objects: 100% (6/6), done.
Delta compression using up to 2 threads
Compressing objects: 100% (4/4), done.
Writing objects: 100% (4/4), 721 bytes | 721.00 KiB/s, done.
Total 4 (delta 2), reused 0 (delta 0), pack-reused 0
To https://gitlab.com/iccn7501450/cicd-demo.git
 2f9f637..efff7d8  master -> master
D77R@vps-4f4e3327:~/cicd-demo$
```

Figure 5.31: `git commit` command triggering *pre-commit* hooks, followed by `git push`

5.5.5.1 Installing a pre-push Hook via pre-commit

Listing 5.1: Installing a pre-push hook with pre-commit

From the project root

```
pre-commit install —hook-type pre-push
```

This command creates (or updates) the `.git/hooks/pre-push` file to call the same checks defined in your `pre-commit` configuration.

```
D77R@vps-4f4e3327:~/cicd-demo$ git push
Username for 'https://gitlab.com': D73L1
Password for 'https://D73L1@gitlab.com':
.git/hooks/pre-push: 16: w#: not found
Running pre-push checks...
[INFO] Scanning for projects...
[INFO]
[INFO] -----< com.lab:microservice-produit >-----
[INFO] Building microservice-produit 0.0.1-SNAPSHOT
[INFO] -----[ jar ]-----
[INFO]
[INFO] --- maven-resources-plugin:3.3.1:resources (default-resources) @ microservice-produit ---
[INFO] Copying 1 resource from src/main/resources to target/classes
[INFO] Copying 1 resource from src/main/resources to target/classes
[INFO]
[INFO] --- maven-compiler-plugin:3.13.0:compile (default-compile) @ microservice-produit ---
[INFO] Nothing to compile - all classes are up to date.
[INFO]
[INFO] --- maven-resources-plugin:3.3.1:testResources (default-testResources) @ microservice-produit ---
[INFO] skip non existing resourceDirectory /home/D77R/cicd-demo/src/test/resources
[INFO]
[INFO] --- maven-compiler-plugin:3.13.0:testCompile (default-testCompile) @ microservice-produit ---
[INFO] Nothing to compile - all classes are up to date.
[INFO]
[INFO] --- maven-surefire-plugin:3.5.2:test (default-test) @ microservice-produit ---
[INFO] Using auto detected provider org.apache.maven.surefire.junitplatform.JUnitPlatformProvider
[INFO]
[INFO] -----
[INFO] T E S T S
[INFO] -----
[INFO] Running com.lab.microserviceproduit.dao.ProductDaoUnitTest
21:23:48.346 [main] INFO org.springframework.test.context.support.AnnotationConfigContextLoaderUtils -- Could not detect default configuration classes for test class [com.lab.microserviceproduit.dao.ProductDaoUnitTest]: ProductDaoUnitTest does not declare any static, non-private, non-final, nested classes annotated with @Configuration.
21:23:48.590 [main] INFO org.springframework.boot.test.context.SpringBootTestContextBootstrapper -- Found @SpringBootConfiguration com.lab.microserviceproduit.MicroserviceProduitApplication for test class com.lab.microserviceproduit.dao.ProductDaoUnitTest
```

Figure 5.32: Example of Maven tests executed before push

```
[INFO] Results:  
[INFO]  
[INFO] Tests run: 11, Failures: 0, Errors: 0, Skipped: 0  
[INFO]  
[INFO] -----  
[INFO] BUILD SUCCESS  
[INFO] -----  
[INFO] Total time: 14.402 s  
[INFO] Finished at: 2024-12-28T21:28:58Z  
[INFO] -----  
All checks passed. Proceeding with push.  
Enumerating objects: 5, done.  
Counting objects: 100% (5/5), done.  
Delta compression using up to 2 threads  
Compressing objects: 100% (3/3), done.  
Writing objects: 100% (3/3), 357 bytes | 357.00 KiB/s, done.  
Total 3 (delta 2), reused 0 (delta 0), pack-reused 0  
To https://gitlab.com/iccn7501450/cicd-demo.git  
 8bb50d7..7de7659 master -> master  
D77R@vps-4f4e3327:~/cicd-demo$
```

Figure 5.33: Output of the pre-push hook: 11 tests executed, no failures

5.5.6 Example Scripts for post-commit and pre-push

```
#!/bin/sh

LATEST_COMMIT=$(git log -1 --pretty=format:'%h - %s (%an)')
BRANCH_NAME=$(git rev-parse --abbrev-ref HEAD)

echo "Sending notification about the commit..."
curl -X POST -H "Content-Type: application/json" \
  -d '{"text":"New commit on branch *'$BRANCH_NAME'*: *'$LATEST_COMMIT'*'}' \
  https://hooks.slack.com/services/your/slack/webhook/url

TAG="tag-$(date +'%Y%m%d%H%M%S')"
git tag $TAG
echo "Tag $TAG created for commit $LATEST_COMMIT"

exit 0
```

Figure 5.34: Example post-commit script (notification and tagging)

This post-commit script notifies a Slack channel with information about the latest commit (BRANCH_NAME, LATEST_COMMIT) and automatically creates a tag.

```
were standard input in the form:
remote="$1"
url="$2"

zero=$(git hash-object --stdin </dev/null | tr '[0-9a-f]' '0')

while read local_ref local_oid remote_ref remote_oid
do
    if test "$local_oid" = "$zero"
    then
        # Handle delete
        :
    else
        if test "$remote_oid" = "$zero"
        then
            # New branch, examine all commits
            range="$local_oid"
        else
            # Update to existing branch, examine new commits
            range="$remote_oid..$local_oid"
        fi

        # Example: Run tests using pytest before allowing the push
        echo "Running pre-push checks..."
        mvn test
        if [ $? -ne 0 ]; then
            echo "Tests failed. Push aborted."
            exit 1
        fi

        # Add any additional custom checks here
        # For example, check for specific patterns in commit messages
        commit=$(git rev-list -n 1 --grep 'FORBIDDEN_PATTERN' "$range")
        if test -n "$commit"
        then
            echo >&2 "Commit with forbidden pattern found in $local_ref. Not pushing."
            exit 1
        fi
    fi
done

echo "All checks passed. Proceeding with push."
exit 0
```

Figure 5.35: Example pre-push script (tests and checks)

This pre-push script runs `mvn test` before the push and can block the push if tests fail or if undesirable patterns (e.g., `FORBIDDEN_PATTERN`) are detected. If no issues are found, the push is allowed.

5.5.7 Strategies to Minimize False Positives

To mitigate false positives generated by automated hooks like `pre-commit`, it is advisable to adopt both technical and methodological approaches:

- **Refinement of validation rules:** Adapt the configurations of tools (Checkstyle, Spotless, OWASP Dependency-Check, etc.) to exclude irrelevant cases. For example, refining analysis rules to ignore certain directories, formats, or patterns known to trigger false alerts. Copy
- **Fine-tuning sensitivity thresholds:** Adjust detection thresholds in static and dynamic analysis tools so that only critical issues are flagged. This technical operation, based on code and behavior analysis, reduces unnecessary interruptions and focuses attention on what matters.
- **Advanced contextual filtering:** Integrate scripts or extensions capable of analyzing the context of changes (branch, author, nature of changes) to distinguish relevant alerts from background noise. This filtering helps target checks without generating false alerts for specific cases.
- **Feedback loop and continuous improvement:** Establish regular feedback with developers to evaluate the relevance of alerts. Feedback allows adjusting tool configurations and refining rules, leading to a gradual reduction in false positives and better alignment with project specifics.

By combining these strategies, the reliability of automated tools is optimized, ensuring they focus on genuine anomalies and enhance code quality and security without hindering the development cycle.

We demonstrate with a small example how to identify and resolve a false positive in Checkstyle by adjusting rules to fit a project's specific conventions.

5.5.8 Step 1: Creating a Java Class with a False Positive

```
public class test {
    private String _name;
    // Espace inutile
}
```

Figure 5.36: Java class `test.java` generating a false positive

In the `test` class, we intentionally introduce:

- A variable `_name`, which violates Checkstyle's default naming rules.
- An unnecessary space after a comment, likely to trigger a warning about unnecessary spaces.

5.5.9 Step 2: Configuring Checkstyle Rules

```
<?xml version="1.0"?>
<!DOCTYPE module PUBLIC
  "-//Checkstyle//DTD Checkstyle Configuration 1.3//EN"
  "https://checkstyle.org/dtds/configuration_1_3.dtd">

<module name="Checker">
  <!-- LineLength module -->
  <module name="LineLength">
    <property name="max" value="120"/>
    <property name="severity" value="warning"/>
  </module>

  <!-- FileTabCharacter module -->
  <module name="FileTabCharacter">
    <property name="severity" value="warning"/>
  </module>

  <module name="LocalVariableName">
    <property name="format" value="^[a-zA-Z][a-zA-Z0-9]*$"/>
  </module>
</module>
```

Figure 5.37: `checkstyle.xml` file with rule configuration

The `checkstyle.xml` file contains the following rules:

- **LineLength**: Limits line length to 120 characters.
- **FileTabCharacter**: Detects the use of tabs.
- **LocalVariableName**: Requires local variable names to follow the format `[a-zA-Z][a-zA-Z0-9]*`.

5.5.10 Step 3: Running Checkstyle on the Class

```
D77R@vps-4f4e3327:~/cicd-demo$ vi checkstyle.xml
D77R@vps-4f4e3327:~/cicd-demo$ checkstyle -c checkstyle.xml test.java
com.puppycrawl.tools.checkstyle.api.CheckstyleException: LocalVariableName is not allowed as a child in Checker
    at com.puppycrawl.tools.checkstyle.Checker.setupChild(Checker.java:502)
    at com.puppycrawl.tools.checkstyle.api.AutomaticBean.configure(AutomaticBean.java:201)
    at com.puppycrawl.tools.checkstyle.Main.runCheckstyle(Main.java:404)
    at com.puppycrawl.tools.checkstyle.Main.runCli(Main.java:331)
    at com.puppycrawl.tools.checkstyle.Main.execute(Main.java:190)
    at com.puppycrawl.tools.checkstyle.Main.main(Main.java:125)
Checkstyle ends with 1 errors.
```

Figure 5.38: Running Checkstyle and the error generated

The following command is used to check the class:

```
checkstyle -c checkstyle.xml test.java
```

Checkstyle generates an error related to the LocalVariableName rule because the variable `_name` violates the format defined in the configuration.

5.5.11 Step 4: Resolving the False Positive

To resolve the issue, we adjust the rule in `checkstyle.xml` to allow variable names starting with an underscore:

```
<module name="LocalVariableName">
<property name="format" value="^[_a-z][a-zA-Z0-9]*\$"/>
</module>
```

5.5.12 Step 5: Validating the Changes

After adjusting the configuration, we rerun Checkstyle:

```
checkstyle -c checkstyle.xml test.java
```

The false positive related to `_name` is resolved, and no errors are reported.

```
D77R@vps-4f4e3327:~/cicd-demo$ vi test.java
D77R@vps-4f4e3327:~/cicd-demo$ checkstyle -c checkstyle.xml test.java
Starting audit...
Audit done.
```

Figure 5.39: Resolution of the False Positive

As a conclusion, suppose a project has:

- A strict Checkstyle rule prohibiting variables starting with an underscore (_).
- Multiple teams allowing this convention as an accepted practice.

Problem: At every commit, the `pre-commit` hook generates unnecessary errors, disrupting developers' workflow.

Implemented Solution:

1. **Adjusting the Checkstyle rule:** Modifying the configuration to include `_name` as a valid format in the `LocalVariableName` rule.
2. **Adding targeted exclusions:** Excluding specific modules, such as test folders or prototypes, to avoid irrelevant alerts.
3. **Conditional validation:** Enabling strict validation only during `pre-push` hooks or on main branches (`main, develop`), reducing the impact on developers' working branches.

5.6 Vulnerability scanning

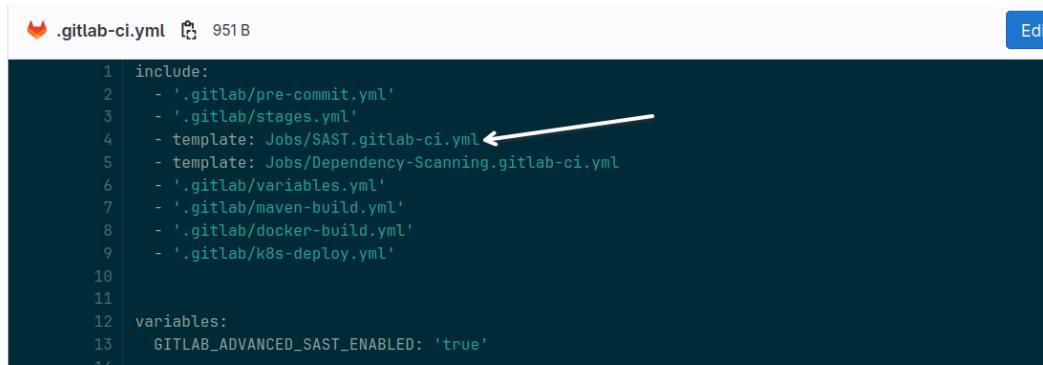
Vulnerability scanning in GitLab is a security feature that is integrated into the CI/CD pipeline to identify and assess potential security vulnerabilities in the codebase and its dependencies. By utilizing automated scanning tools, GitLab can analyze the project for known vulnerabilities in libraries, containers, and configurations during the build process. GitLab provides many types of vulnerability scanning, notably:

- **Static Application Security Testing (SAST):** This scanning method analyzes source code for security vulnerabilities without executing the application. It helps identify issues such as code injection and insecure coding practices early in the development process.
- **Software Composition Analysis (SCA):** SCA examines the project's dependencies and libraries for known vulnerabilities. It ensures that third-party components do not introduce security risks into the application.
- **Dynamic Application Security Testing (DAST):** DAST scans the running application to identify vulnerabilities that can be exploited during runtime, such as cross-site scripting (XSS) and SQL injection.

- **Container Scanning:** This type of scanning assesses container images for known vulnerabilities before deployment. It checks both the base images and any installed packages within the container for security issues.

5.6.1 Static Application Security Testing

Lucky for SAST, GitLab comes with pre-made templates that can be easily included in our CI/CD pipeline as shown in figure 5.47. Note this template contains a job that's part of a stage named test, thus we also added a stage test as the first stage in the file 'stages.yml'



```

.gitlab-ci.yml  951B
Edit

1 include:
2   - '.gitlab/pre-commit.yml'
3   - '.gitlab/stages.yml'
4   - template: Jobs/SAST.gitlab-ci.yml ←
5   - template: Jobs/Dependency-Scanning.gitlab-ci.yml
6   - '.gitlab/variables.yml'
7   - '.gitlab/maven-build.yml'
8   - '.gitlab/docker-build.yml'
9   - '.gitlab/k8s-deploy.yml'
10
11
12 variables:
13   GITLAB_ADVANCED_SAST_ENABLED: 'true'
14

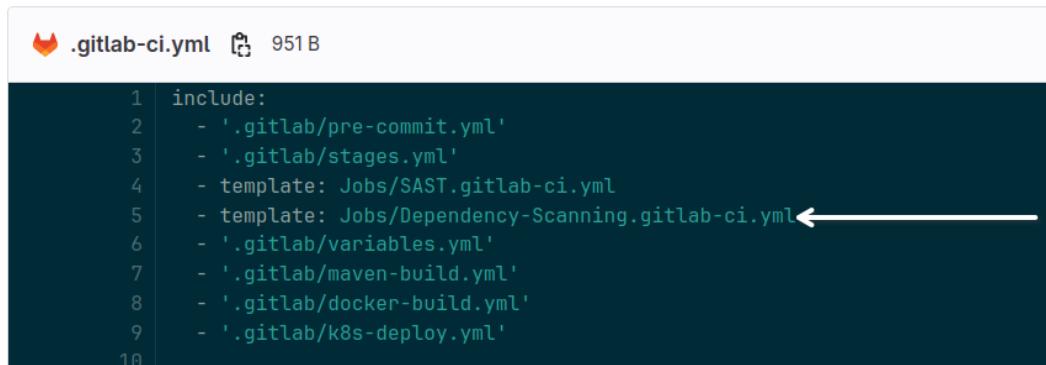
```

Figure 5.40

The scan results are saved in an artifact that we can download.

5.6.2 Software Composition Analysis

Similarly to enabling SAST, GitLab provides a template that contains a job that does software composition analysis. And we only need to include it.



```

.gitlab-ci.yml  951B
Edit

1 include:
2   - '.gitlab/pre-commit.yml'
3   - '.gitlab/stages.yml'
4   - template: Jobs/SAST.gitlab-ci.yml
5   - template: Jobs/Dependency-Scanning.gitlab-ci.yml ←
6   - '.gitlab/variables.yml'
7   - '.gitlab/maven-build.yml'
8   - '.gitlab/docker-build.yml'
9   - '.gitlab/k8s-deploy.yml'
10

```

Figure 5.41

The scan results are saved in an artifact that we can download (figure 5.42). Or in what is coming next we can visualize the discovered vulnerabilities in GitLab Security Dashboard.



Figure 5.42

5.6.3 Container Security

The objective for Container Security is to scan the Docker image for vulnerabilities using **Docker Scout** and ensure no critical vulnerabilities are present before deployment.

Implementation Steps

1. Define the **docker-scout** Job in `.gitlab-ci.yml`: The `docker-scout` job is added to the CI/CD pipeline to scan the Docker image after it is built. The job performs the following steps:

- Install Docker Scout using `curl`.
- Log in to Docker Hub using the provided credentials.
- Scan the Docker image for vulnerabilities using Docker Scout.

`docker-scout:`

```

stage: security_scan
image: docker:latest
services:
  - docker:dind
variables:
  DOCKER_IMAGE: "${DOCKER_LOGIN}/${IMAGE_NAME}"
  DOCKER_HUB_USERNAME: "$DOCKER_LOGIN"
  DOCKER_HUB_TOKEN: "$DOCKER_TOKEN"
script:
  - echo "Installing Docker Scout..."
  - apk add --no-cache curl tar file
  - echo "Logging in to Docker Hub..."
```

```

- echo "$DOCKER_HUB_TOKEN" | docker login -u "$DOCKER_HUB_USERNAME" --password-stdin
- echo "Scanning Docker image for vulnerabilities..."
- curl -fsSL https://github.com/docker/scout-cli/releases/download/v1.16.1/docker-scout
- tar -xzf docker-scout.tar.gz -C /usr/local/bin docker-scout
- chmod +x /usr/local/bin/docker-scout
- /usr/local/bin/docker-scout cves $DOCKER_IMAGE

# container security
docker-scout:
  stage: container_scan
  image: docker:latest
  services:
    - docker:dind
  variables:
    DOCKER_IMAGE: "${DOCKER_LOGIN}/${IMAGE_NAME}"
    DOCKER_HUB_USERNAME: "$DOCKER_LOGIN"
    DOCKER_HUB_TOKEN: "$DOCKER_TOKEN"
  script:
    - echo "Installing Docker Scout..."
    - apk add --no-cache curl tar file # Install curl, tar, and file
    - echo "Checking system architecture..."
    - uname -m
    - echo "Checking Docker version..."
    - docker --version
    - echo "Logging in to Docker Hub..."
    - echo "$DOCKER_HUB_TOKEN" | docker login -u "$DOCKER_HUB_USERNAME" --password-stdin
    - echo "Verifying Docker login..."
    - docker info
    - echo "Downloading Docker Scout version 1.16.1 for AMD64..."
    - curl -fsSL https://github.com/docker/scout-cli/releases/download/v1.16.1/docker-scout_1.16.1_linux_amd64.tar.gz -o docker-scout.tar.gz
    - echo "Verifying the downloaded file..."
    - file docker-scout.tar.gz
    - echo "Listing contents of the archive..."
    - tar -tf docker-scout.tar.gz
    - echo "Extracting the binary..."
    - tar -xzf docker-scout.tar.gz -C /usr/local/bin docker-scout
    - chmod +x /usr/local/bin/docker-scout
    - echo "Checking the binary..."
    - ls -l /usr/local/bin/docker-scout
    - echo "Testing Docker Scout binary..."
    - /usr/local/bin/docker-scout --help
    - echo "Testing Docker Scout with a simple command..."
    - /usr/local/bin/docker-scout cves hello-world
    - echo "Scanning Docker image for vulnerabilities..."
    - /usr/local/bin/docker-scout cves $DOCKER_IMAGE
  allow_failure: false

```

Figure 5.43: configuration for Docker Scout inside a container

2. Results: The Docker image is successfully scanned, and no critical vulnerabilities are detected. The pipeline continues without interruption, confirming the image meets security standards.

```

714   0C  0H  0M   1L util-linux 2.36.1-8+deb11u1
715 pkg:deb/debian/util-linux@2.36.1-8%2Bdeb11u1?os_distro=bullseye&os_name=debian&os_version=11
716   x LOW CVE-2024-28085
717     https://scout.docker.com/v/CVE-2024-28085?s=debian&n=util-linux&ns=debian&t=deb&osn=debian&osv=11&vr=%3C2.36.1-8%2Bdeb11u2
718       Affected range : <2.36.1-8+deb11u2
719       Fixed version  : 2.36.1-8+deb11u2
720
721   0C  0H  0M   1L gcc-9 9.3.0-22
722 pkg:deb/debian/gcc-9@9.3.0-22?os_distro=bullseye&os_name=debian&os_version=11
723   x LOW CVE-2023-4039
724     https://scout.docker.com/v/CVE-2023-4039?s=debian&n=gcc-9&ns=debian&t=deb&osn=debian&osv=11&vr=%3E%3D9.3.0-22
725       Affected range : >9.3.0-22
726       Fixed version  : not fixed
727
728 90 vulnerabilities found in 28 packages
729   CRITICAL    12
730   HIGH        28
731   MEDIUM      25
732   LOW         21
733   UNSPECIFIED 4
734 What's next:
735   View base image update recommendations → docker scout recommendations projdevopsiccn2/mcommerce:lates
t
▼ 736 Cleaning up project directory and file based variables
  00:01
737 Job succeeded

```

Figure 5.44: docker scout job results

5.6.4 Dynamic Application Security Testing

1. Define the `dast` Job in `.gitlab-ci.yml`: The `dast` job is added to the pipeline to scan the running application for vulnerabilities. The job performs the following steps:

- Deploy the application using Docker.
- Run the DAST scan on the application's endpoints.
- Save the scan results as artifacts for further analysis.

```

dast:
  stage: security_scan
  services:
    - name: ${DOCKER_LOGIN}/${IMAGE_NAME}
      alias: app
  variables:
    DAST_WEBSITE: "http://app:8080"
  script:
    - echo "Starting DAST scan..."
    - /analyze

```

```
artifacts:
```

```
paths:
```

- dast_artifacts/

```

1  image: docker:latest
2  services:
3    - docker:dind
4
5  include:
6    - '.gitlab/pre-commit.yml'
7    - '.gitlab/stages.yml'
8    - template: Jobs/SAST.gitlab-ci.yml
9    - template: Security/DAST.gitlab-ci.yml # Correct DAST template
10   - template: Jobs/Dependency-Scanning.gitlab-ci.yml
11   - '.gitlab/variables.yml'
12   - '.gitlab/maven-build.yml'
13   - '.gitlab/docker-build.yml'
14   - '.gitlab/k8s-deploy.yml'
15   - '.gitlab/delete_old_pipeline.yml'
16
17 variables:
18   GITLAB_ADVANCED_SAST_ENABLED: 'true'
19   DOCKER_LOGIN: "projdevopsiccn2"
20   IMAGE_NAME: "mcommerce"
21   DAST_WEBSITE: "http://app:8080" # URL for DAST scan
22

```

Figure 5.45: dast test configuration -1-

```

# Override the DAST job stage
dast:
  stage: dast
  services:
    - name: ${DOCKER_LOGIN}/${IMAGE_NAME} # Run the application as a service
      alias: app
  variables:
    DAST_WEBSITE: "http://app:8080" # URL for DAST scan
  script:
    - echo "Starting DAST scan..."
      - /analyze
  artifacts:
    paths:
      - dast_artifacts/
  allow_failure: false

```

Figure 5.46: dast test configuration -2-

2. Results: The DAST scan is executed successfully, and no critical vulnerabilities are detected in the application. The pipeline continues without interruption, confirming the application meets security requirements.

Conclusion

All tests, including **SAST**, **DAST**, **Container Security**, and **Pre-Commit Hooks**, passed successfully. The Docker image is free from critical vulnerabilities, the application is secure against runtime threats, and the code meets quality standards. The Product Microservice is now ready for deployment.

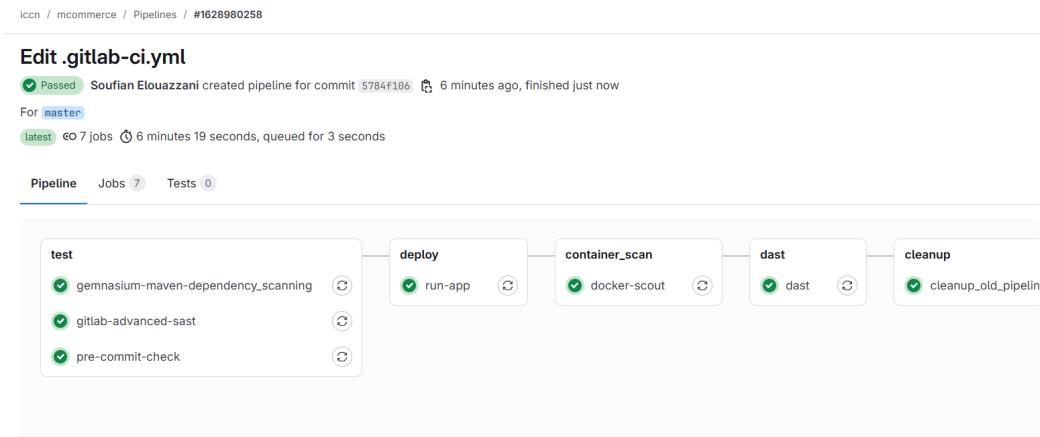


Figure 5.47: pipeline for security testing gone Well

5.7 Vulnerability management and notification

To visualize the detected vulnerabilities in the different scans in our CI pipeline we need to use a vulnerability management platform. For our case we chose GitLab Security dashboard because it's already part of the tooling we have been using in this project and easy to work with.

5.7.1 Vulnerability management

After enabling scanning such as DAST or SCA, scan results are automatically shown in the vulnerability report.

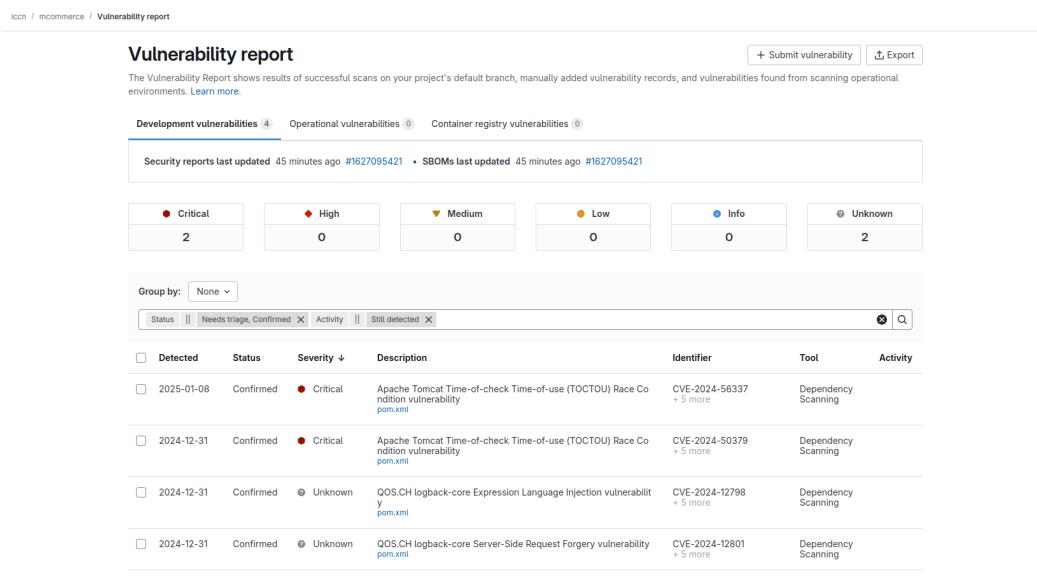


Figure 5.48

5.7.2 Configuring notification using Slack

GitLab support sending event notifications such as finding a new vulnerability to Slack. To configure this. First we will create a Slack Workspace and a channel for security notifications.

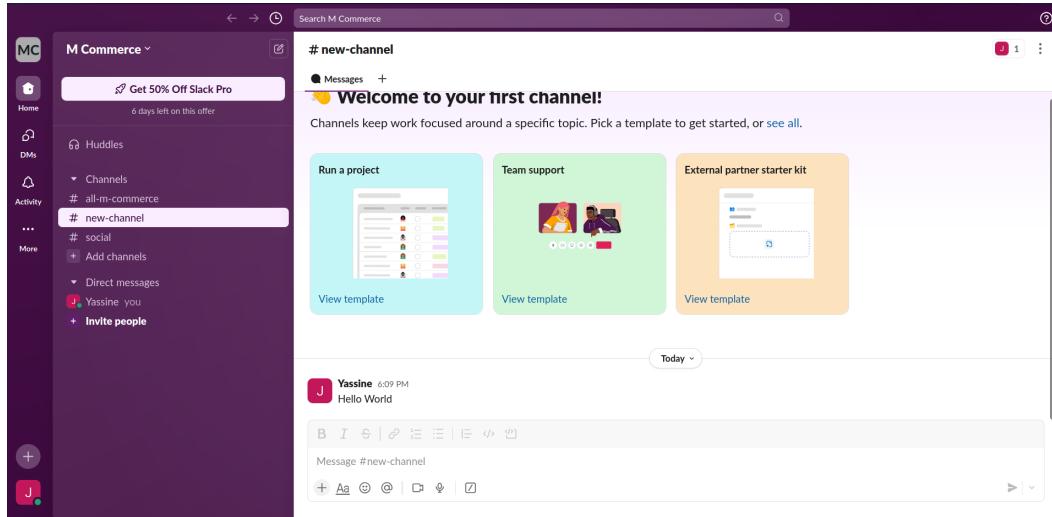


Figure 5.49: Creating a Slack workplace

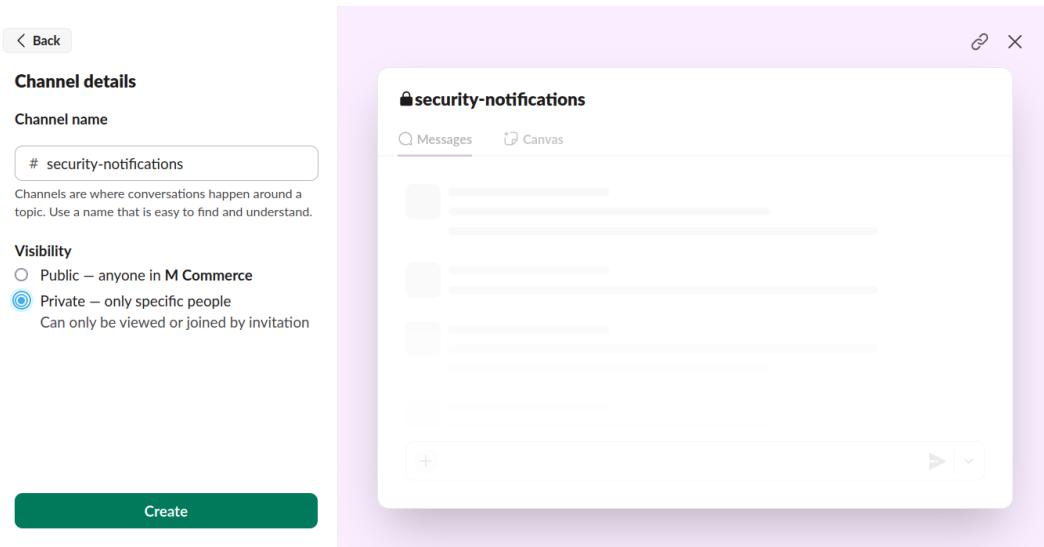


Figure 5.50: Creating a channel for vulnerabilities notifications

Then we connect Slack with our GitLab repository.

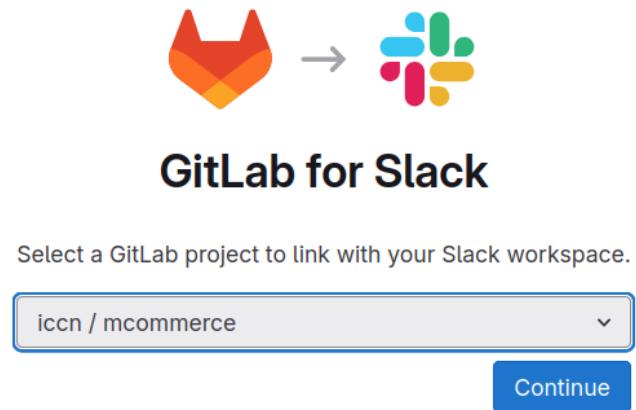


Figure 5.51

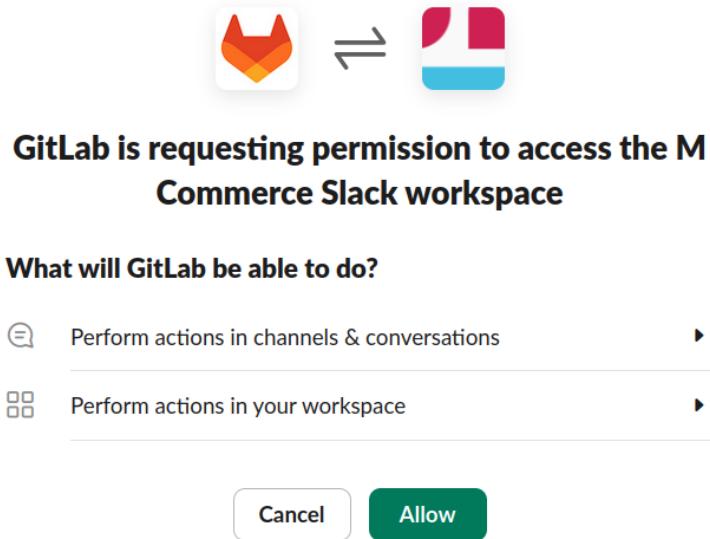


Figure 5.52

We then add GitLab to our security-notifications channel.

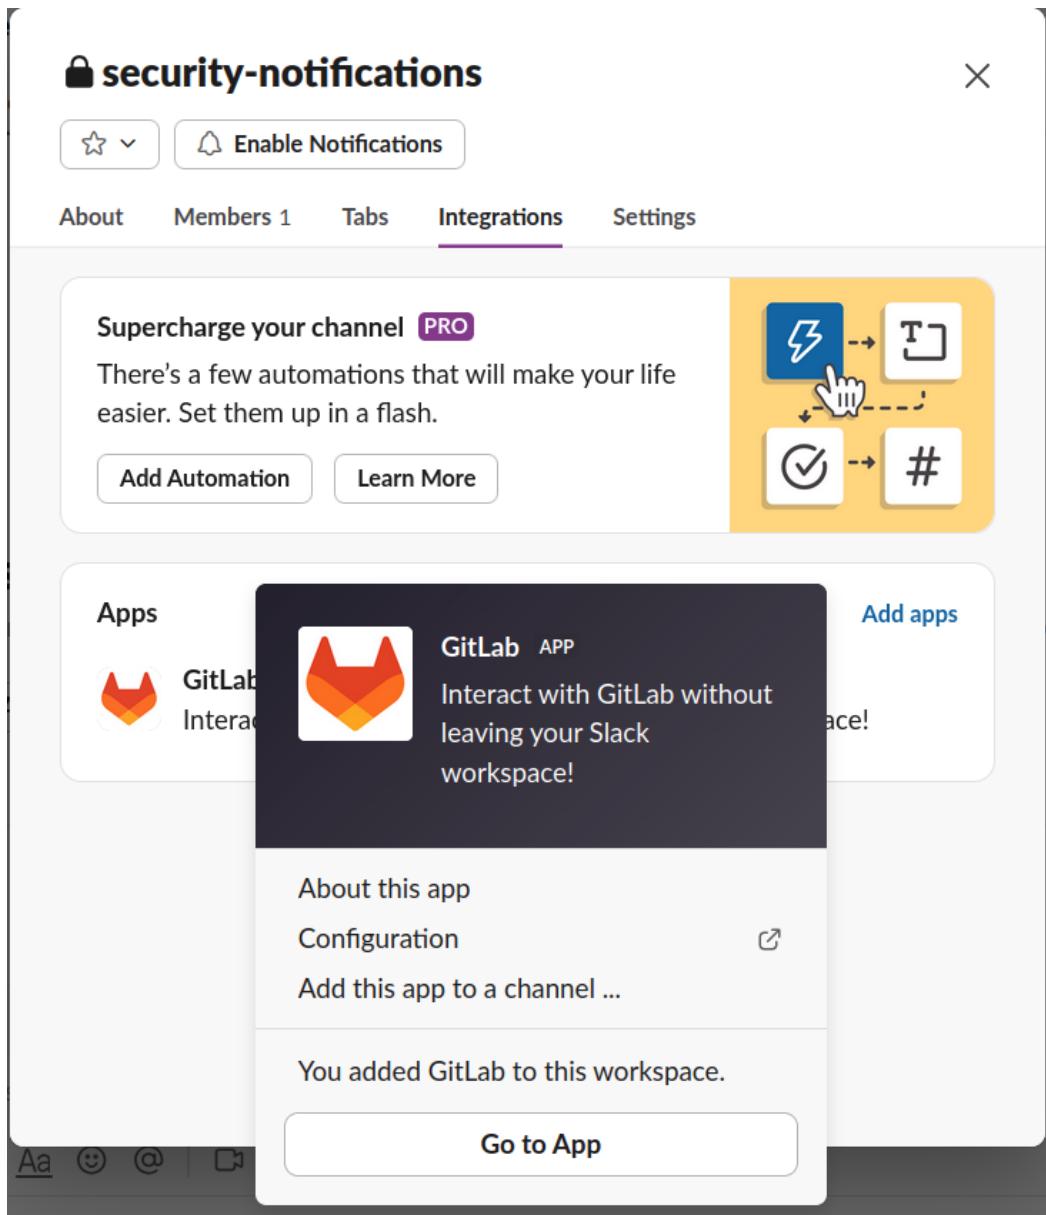


Figure 5.53

And finally we configure GitLab to send new vulnerabilities notifications to the channel security-notifications.

A comment is added

An internal note or comment on a confidential issue is added

A tag is pushed to the repository or removed

A pipeline status changes

A wiki page is created or updated

A deployment is started or finished

An incident is created, closed, or reopened

```
#security-notifications
```

A new, unique vulnerability is recorded (available only in GitLab Ultimate)

```
#security-notifications
```

A new, unique alert is recorded

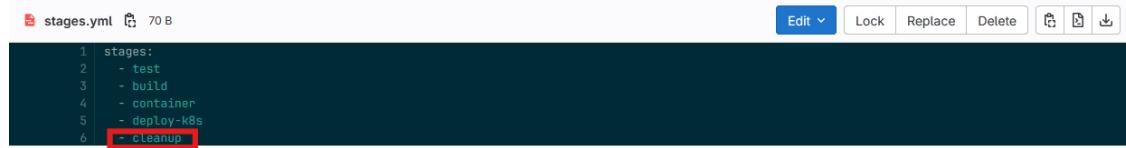
```
#security-notifications
```

Notification settings

Figure 5.54

5.8 Automated Pipeline Cleaning

First, we added the `cleanup` stage to the list of stages in `stages.yml` as shown below:



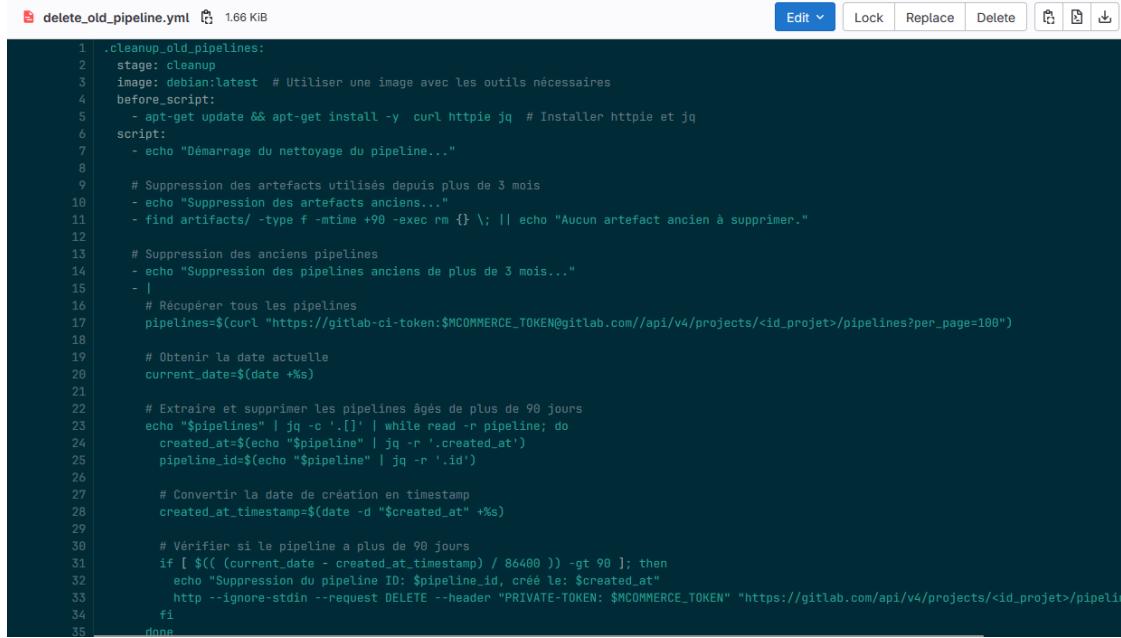
```
stages.yml 70 B
Edit Lock Replace Delete
1 stages:
2   - test
3   - build
4   - container
5   - deploy-k8s
6   - cleanup
```

Figure 5.55: Stages.yml updated with the cleanup stage.

Then, we created the `delete old pipeline.yaml` file in which we included the instructions for running our stage:

To ensure the pipeline supports it, we included the stage in the `.gitlab-ci.yml` pipeline file and made an extension of the job file:

Finally, we checked the pipeline to verify that our job is working well. We can observe that the pipeline is executing correctly:

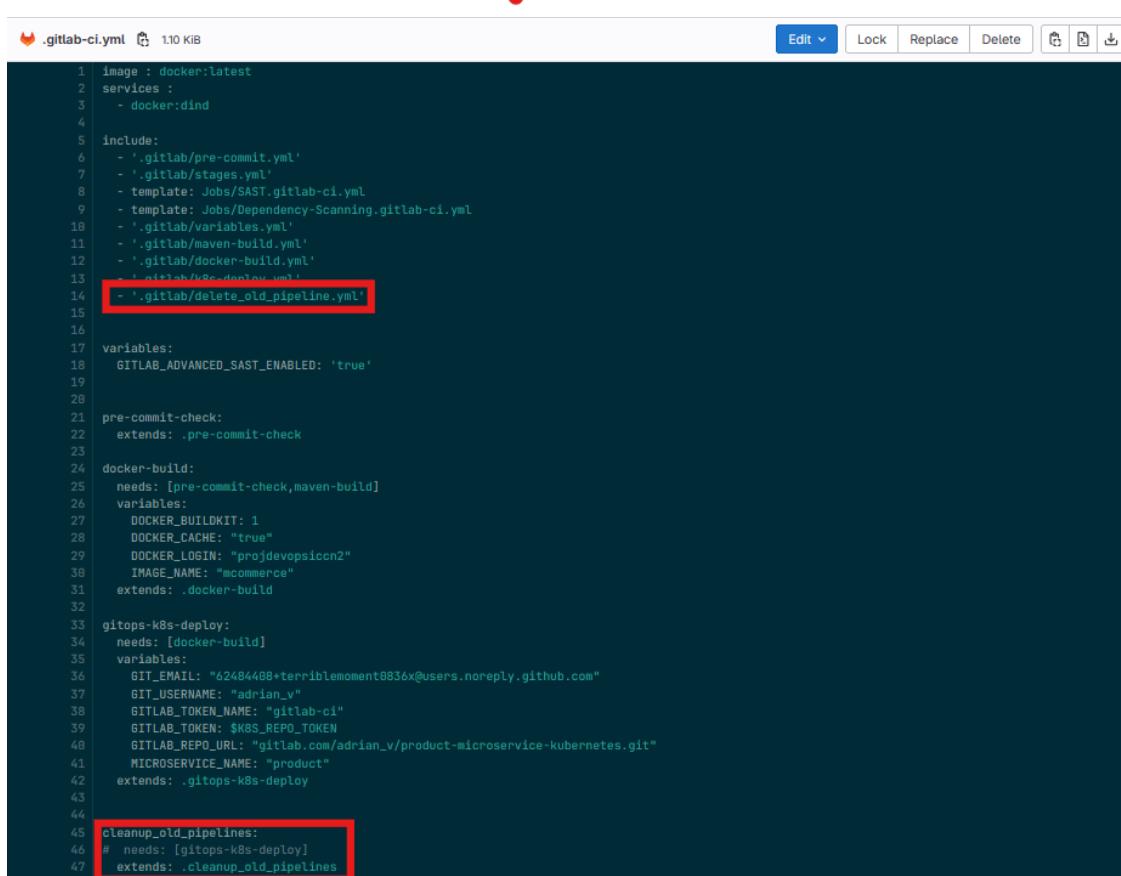


```

1 .cleanup_old_pipelines:
2   stage: cleanup
3   image: debian:latest # Utiliser une image avec les outils nécessaires
4   before_script:
5     - apt-get update && apt-get install -y curl httpie jq # Installer httpie et jq
6   script:
7     - echo "Démarrage du nettoyage du pipeline..."
8
9     # Suppression des artefacts utilisés depuis plus de 3 mois
10    - echo "Suppression des artefacts anciens..."
11    - find artifacts/ -type f -mtime +90 -exec rm {} \; || echo "Aucun artefact ancien à supprimer."
12
13    # Suppression des anciens pipelines
14    - echo "Suppression des pipelines anciens de plus de 3 mois..."
15    - |
16      # Récupérer tous les pipelines
17      pipelines=$(curl "https://gitlab-ci-token:$MCOMMERCE_TOKEN@gitlab.com/api/v4/projects/<id_projet>/pipelines?per_page=100")
18
19      # Obtenir la date actuelle
20      current_date=$(date +\%s)
21
22      # Extraire et supprimer les pipelines âgés de plus de 90 jours
23      echo "$pipelines" | jq -c '.[]' | while read -r pipeline; do
24        created_at=$(echo "$pipeline" | jq -r '.created_at')
25        pipeline_id=$(echo "$pipeline" | jq -r '.id')
26
27        # Convertir la date de création en timestamp
28        created_at_timestamp=$(date -d "$created_at" +\%s)
29
30        # Vérifier si le pipeline a plus de 90 jours
31        if [ $(( (current_date - created_at_timestamp) / 86400 )) -gt 90 ]; then
32          echo "Suppression du pipeline ID: $pipeline_id, créé le: $created_at"
33          http --ignore-stdin --request DELETE --header "PRIVATE-TOKEN: $MCOMMERCE_TOKEN" "https://gitlab.com/api/v4/projects/<id_projet>/pipelines/$pipeline_id"
34        fi
35      done

```

Figure 5.56: Instructions for the cleanup stage.



```

1 image : docker:latest
2 services :
3   - docker:dind
4
5 include:
6   - '.gitlab/pre-commit.yml'
7   - '.gitlab/stages.yml'
8   - template: Jobs/SAST.gitlab-ci.yml
9   - template: Jobs/Dependency-Scanning.gitlab-ci.yml
10  - '.gitlab/variables.yml'
11  - '.gitlab/maven-build.yml'
12  - '.gitlab/docker-build.yml'
13  - '.gitlab/docker-deploy.yml'
14  - '.gitlab/delete_old_pipeline.yml'
15
16
17 variables:
18   GITLAB_ADVANCED_SAST_ENABLED: 'true'
19
20
21 pre-commit-check:
22   extends: .pre-commit-check
23
24 docker-build:
25   needs: [pre-commit-check,maven-build]
26   variables:
27     DOCKER_BUILDKIT: 1
28     DOCKER_CACHE: "true"
29     DOCKER_LOGIN: "projdevopsiccon2"
30     IMAGE_NAME: "mcommerce"
31   extends: .docker-build
32
33 gitops-k8s-deploy:
34   needs: [docker-build]
35   variables:
36     GIT_EMAIL: "62484408+terribllement0036x@users.noreply.github.com"
37     GIT_USERNAME: "adrian_v"
38     GITLAB_TOKEN_NAME: "gitlab-ci"
39     GITLAB_TOKEN: "$K8S_REPO_TOKEN"
40     GITLAB_REPO_URL: "gitlab.com/adrian_v/product-microservice-kubernetes.git"
41     MICROSERVICE_NAME: "product"
42   extends: .gitops-k8s-deploy
43
44
45 cleanup_old_pipelines:
46   # needs: [gitops-k8s-deploy]
47   extends: .cleanup_old_pipeline

```

Figure 5.57: Extension of the pipeline job file.



Figure 5.58: Pipeline execution verification.

Chapter 6

Observability and Monitoring

6.1 infrastructure Monitoring

6.1.1 Installing Prometheus on Kubernetes

```
D77R@vps-4f4e3327:/home/marouane$ helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
"prometheus-community" has been added to your repositories
D77R@vps-4f4e3327:/home/marouane$ helm repo add stable https://charts.helm.sh/stable
"stable" has been added to your repositories
D77R@vps-4f4e3327:/home/marouane$ helm repo update
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "prometheus-community" chart repository
...Successfully got an update from the "stable" chart repository
Update Complete. #Happy Helm-ing!
D77R@vps-4f4e3327:/home/marouane$
```

Figure 6.1: Add prometheus chart to Helm repo

```
D77R@vps-4f4e3327:/home/marouane$ helm install prometheus prometheus-community/kube-prometheus-stack --timeout 10m
NAME: prometheus
LAST DEPLOYED: Mon Jan 13 15:09:16 2025
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
kube-prometheus-stack has been installed. Check its status by running:
  kubectl --namespace default get pods -l "release=prometheus"
Visit https://github.com/prometheus-operator/kube-prometheus for instructions on how to create & configure Alertmanager
and Prometheus instances using the Operator.
D77R@vps-4f4e3327:/home/marouane$
```

Figure 6.2: Installing Prometheus on Kubernetes

```
D77R@vps-4f4e3327:/home/marouane$ kubectl get all
NAME                                         READY   STATUS    RESTARTS   AGE
pod/prometheus-grafana-37094ff0d5-ctvld7   3/3     Running   0          8m55s
pod/prometheus-kube-prometheus-operator-64cc85bb7f-mxzl2   0/1     CrashLoopBackOff   1          8m55s
pod/prometheus-kube-state-metrics-6f5574c548-zwitt      0/1     Running   3 (32s ago)   8m55s
pod/prometheus-node-exporter-vd7wn           1/1     Running   2 (88s ago)   8m55s

NAME                           TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/kubernetes             ClusterIP   <none>        <none>       443/TCP   23d
service/prometheus-grafana     ClusterIP   10.96.8.1    <none>       80/TCP    8m55s
service/prometheus-kube-prometheus-alertmanager   ClusterIP   10.10.98.196  <none>       9093/TCP,8080/TCP   8m55s
service/prometheus-kube-prometheus-operator   ClusterIP   10.99.25.234  <none>       443/TCP   8m55s
service/prometheus-kube-prometheus   ClusterIP   10.10.98.196  <none>       9093/TCP,8080/TCP   8m55s
service/prometheus-kube-state-metrics   ClusterIP   10.100.107.40  <none>       8088/TCP   8m55s
service/prometheus-node-exporter      ClusterIP   10.104.22.197 <none>       9100/TCP   8m55s

NAME              DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
daemonset.apps/prometheus-node-exporter   1         1         1         1           1           kubernetes.io/os:linux   8m55s

NAME                                         READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/prometheus-grafana     1/1     1           1           8m55s
deployment.apps/prometheus-kube-prometheus-operator 0/1     1           0           8m55s
deployment.apps/prometheus-kube-state-metrics 0/1     1           0           8m55s

NAME                                         DESIRED   CURRENT   READY   AGE
replicaset.apps/prometheus-grafana-578946f5d5   1         1         1         8m55s
replicaset.apps/prometheus-kube-prometheus-operator-64cc85bb7f  1         1         0         8m55s
replicaset.apps/prometheus-kube-state-metrics-6f5574c548  1         1         0         8m55s
D77R@vps-4f4e3327:/home/marouane$
```

Figure 6.3: Creation of prometheus components : pods ,service and daemon-set ...

NAME	DATA	AGE
kube-root-ca.crt	1	23d
prometheus-grafana	1	9m16s
prometheus-grafana-config-dashboards	1	9m16s
prometheus-kube-prometheus-alertmanager-overview	1	9m16s
prometheus-kube-prometheus-apiserver	1	9m16s
prometheus-kube-prometheus-cluster-total	1	9m16s
prometheus-kube-prometheus-controller-manager	1	9m16s
prometheus-kube-prometheus-etcd	1	9m16s
prometheus-kube-prometheus-grafana-datasource	1	9m16s
prometheus-kube-prometheus-grafana-overview	1	9m16s
prometheus-kube-prometheus-k8s-coredns	1	9m16s
prometheus-kube-prometheus-k8s-resources-cluster	1	9m16s
prometheus-kube-prometheus-k8s-resources-multicluster	1	9m16s
prometheus-kube-prometheus-k8s-resources-namespace	1	9m16s
prometheus-kube-prometheus-k8s-resources-node	1	9m16s
prometheus-kube-prometheus-k8s-resources-pod	1	9m16s
prometheus-kube-prometheus-k8s-resources-workload	1	9m16s
prometheus-kube-prometheus-k8s-resources-workloads-namespace	1	9m16s
prometheus-kube-prometheus-kubelet	1	9m16s
prometheus-kube-prometheus-namespace-by-pod	1	9m16s
prometheus-kube-prometheus-namespace-by-workload	1	9m16s
prometheus-kube-prometheus-node-cluster-rsrc-use	1	9m16s
prometheus-kube-prometheus-node-rsrc-use	1	9m16s
prometheus-kube-prometheus-nodes	1	9m16s
prometheus-kube-prometheus-nodes-aix	1	9m16s
prometheus-kube-prometheus-nodes-darwin	1	9m16s
prometheus-kube-prometheus-persistentvolumesusage	1	9m16s
prometheus-kube-prometheus-pod-total	1	9m16s
prometheus-kube-prometheus-prometheus	1	9m16s
prometheus-kube-prometheus-proxy	1	9m16s
prometheus-kube-prometheus-scheduler	1	9m16s
prometheus-kube-prometheus-workload-total	1	9m16s

Figure 6.4: Prometheus ConfigMap

NAME	TYPE	DATA	AGE
alertmanager-prometheus-kube-prometheus-alertmanager	Opaque	1	10m
prometheus-grafana	Opaque	3	10m
prometheus-kube-prometheus-admission	Opaque	3	10m
sh.helm.release.v1.prometheus.v1	helm.sh/release.v1	1	14m

Figure 6.5: Secrets

NAME	CREATED AT
alertmanagerconfigs.monitoring.coreos.com	2025-01-13T15:06:12Z
alertmanagers.monitoring.coreos.com	2025-01-13T15:06:12Z
applications.argojob.io	2024-12-21T14:26:22Z
applicationsets.argojob.io	2024-12-21T14:26:23Z
appprojects.argojob.io	2024-12-21T14:26:23Z
podmonitors.monitoring.coreos.com	2025-01-13T15:06:13Z
probes.monitoring.coreos.com	2025-01-13T15:06:13Z
prometheusagents.monitoring.coreos.com	2025-01-13T15:06:13Z
prometheuses.monitoring.coreos.com	2025-01-13T15:06:14Z
prometheusrules.monitoring.coreos.com	2025-01-13T15:06:15Z
scrapeconfigs.monitoring.coreos.com	2025-01-13T15:06:15Z
servicemonitors.monitoring.coreos.com	2025-01-13T15:06:16Z
thanosrulers.monitoring.coreos.com	2025-01-13T15:06:17Z

Figure 6.6: Custom Resource Definition CRD

6.1.2 Access Prometheus UI

NAME	READY	STATUS	RESTARTS	AGE
prometheus-grafana-578946f5d5-qtnb7	3/3	Running	3 (5m11s ago)	11m
prometheus-kube-prometheus-operator-64cc05eb7f-mxlzl	0/1	CrashLoopBackoff	5 (33s ago)	11m
prometheus-kube-state-metrics-6f5574c548-jvlft	0/1	CrashLoopBackoff	4 (44s ago)	11m
prometheus-prometheus-node-exporter-vd7wm	1/1	Running	6	11m

Figure 6.7: Pods

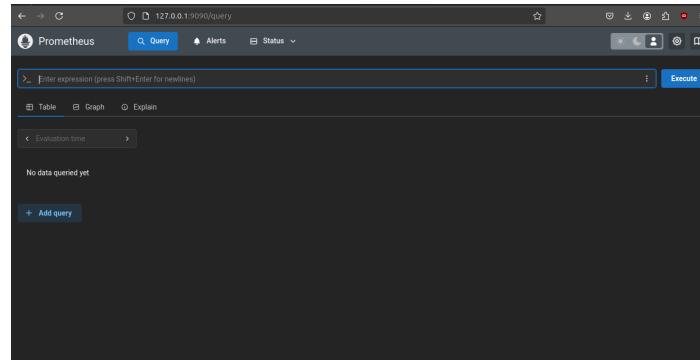


Figure 6.8: Prometheus WEB UI

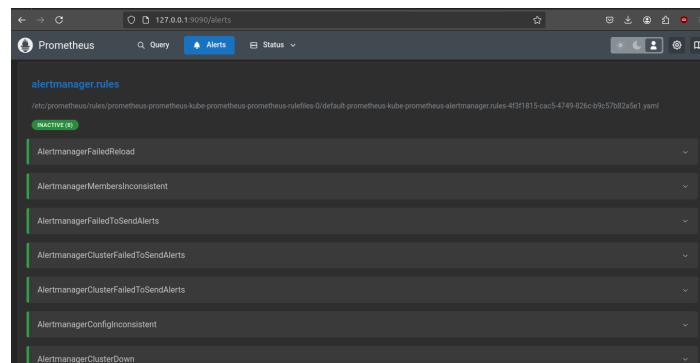


Figure 6.9: Prometheus Alerts

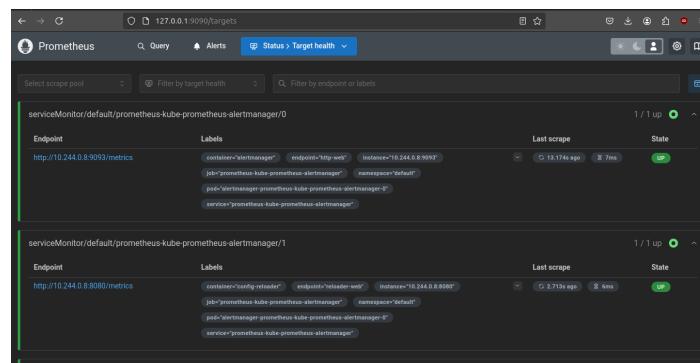


Figure 6.10: Prometheus Taragets

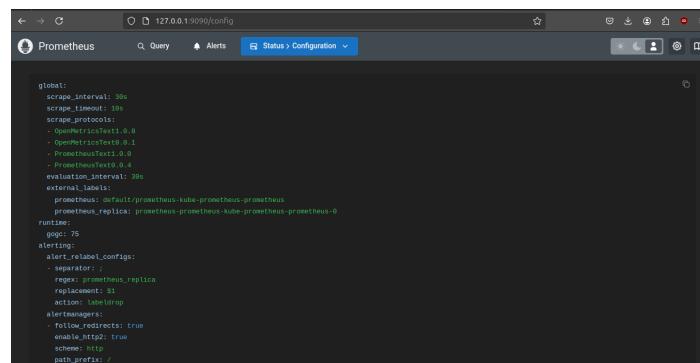


Figure 6.11: Prometheus yaml file for configurations

6.1.3 Graphana

```
D77R@vps-4f4e3327:~$ kubectl get pod
NAME                                         READY   STATUS    RESTARTS   AGE
alertmanager-prometheus-kube-prometheus-alertmanager-0   2/2     Running   0          38m
grafana-69d855495d-lhm6k                      0/1     Running   0          18s
prometheus-grafana-578946f5d5-jm8pw           3/3     Running   0          38m
prometheus-kube-prometheus-operator-6cd5954785-jw7n6   1/1     Running   0          38m
prometheus-kube-state-metrics-6f5574c548-kvrq9      1/1     Running   0          38m
prometheus-prometheus-kube-prometheus-prometheus-0   2/2     Running   0          38m
prometheus-prometheus-node-exporter-nkfv2          1/1     Running   0          38m
D77R@vps-4f4e3327:~$
```

Figure 6.12: Graphana Pod

```
D77R@vps-4f4e3327:/home/marouane$ kubectl port-forward prometheus-grafana-578946f5d5-jm8pw 3001
Forwarding from 127.0.0.1:3001 -> 3001
Forwarding from [::1]:3001 -> 3001
```

Figure 6.13: Access the Grafana application in a cluster using Port Forwarding

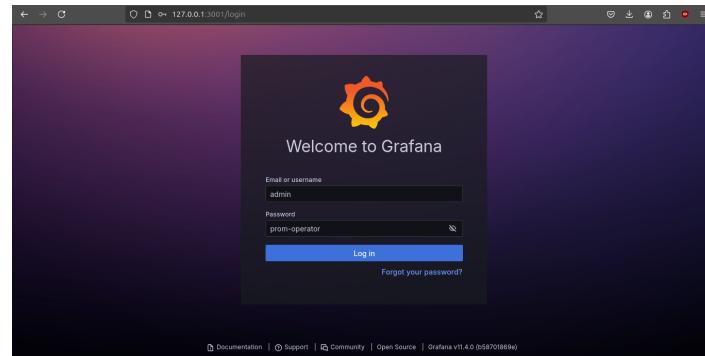


Figure 6.14: Graphana login

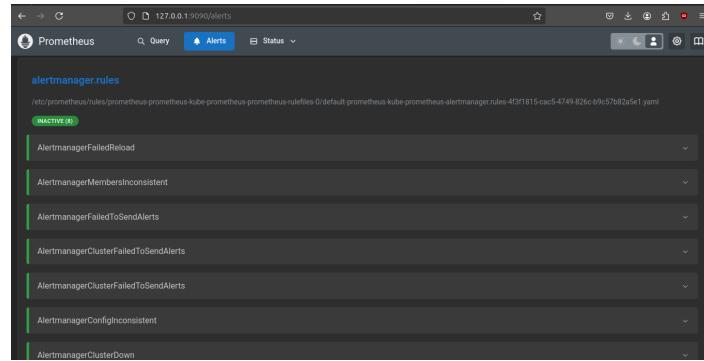


Figure 6.15: alerts

Graphana also offers a user-freindly interface for mesuring system metrics suj as : RAM , CPU , network

...



Figure 6.16: CPU



Figure 6.17: Memory

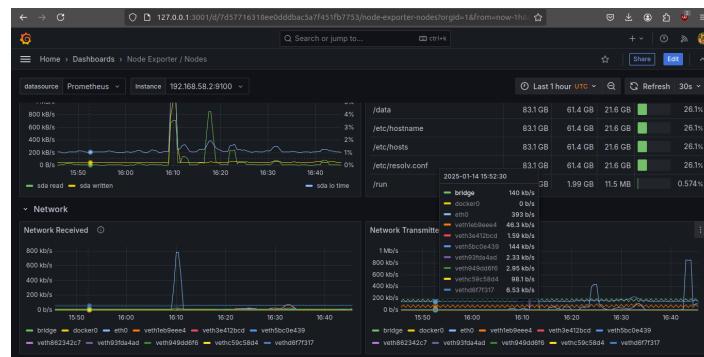


Figure 6.18: Network

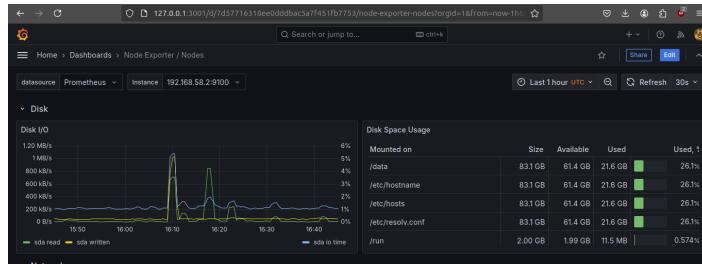


Figure 6.19: Disk usage

6.2 Implementation of Application Monitoring for the Produit Microservice

To ensure the reliability and performance of the **Produit** microservice, we integrated **Prometheus** and **Grafana** into our GitLab CI/CD pipeline. This setup enables real-time monitoring and visualization of the microservice's behavior.

6.2.1 Why Prometheus and Grafana?

- **Prometheus:** An open-source monitoring tool that collects metrics from applications and provides a powerful query language (**PromQL**) for analysis.
- **Grafana:** A widely-used visualization platform that integrates with Prometheus to provide detailed insights through custom dashboards and alerts.

6.2.2 Implementation Steps

6.2.2.1 Prometheus Setup

1. Added Prometheus to the GitLab pipeline as a service using the `prom/prometheus:latest` image.
2. Created a `prometheus.yml` file to scrape metrics from the microservice's endpoint (`/metrics`).
3. Integrated Prometheus into the Maven project by exposing metrics using the Prometheus Java Client library.

6.2.2.2 Grafana Setup

1. Added Grafana to the GitLab pipeline as a service using the `grafana/grafana:latest` image.

2. Configured Grafana to use Prometheus as a data source.

6.2.3 Screenshots and Observations

Below are screenshots capturing key aspects of the implementation:

- Prometheus scraping metrics from the **Product** microservice (Figure 6.20).
- Grafana dashboard visualizing the metrics (Figure 6.21).

```

111  prometheus:
112    stage: monitoring
113    image: prom/prometheus:latest
114    script:
115      - echo "Prometheus service is running."
116    artifacts:
117      paths:
118        - prometheus.yml
119    variables:
120      PROMETHEUS_CONFIG_PATH: prometheus.yml
121    allow_failure: false

```

Figure 6.20: Prometheus configuration and metrics scraping.

```

123  grafana:
124    stage: monitoring
125    image: grafana/grafana:latest
126    script:
127      - echo "Grafana service is running."
128      - echo "Grafana is accessible at http://localhost:3000 with default credentials (admin/admin)."
129    allow_failure: false

```

Figure 6.21: Grafana dashboard visualizing metrics.

6.2.4 Conclusion

The integration of Prometheus and Grafana into the GitLab CI/CD pipeline provided real-time visibility into the **Product** microservice's performance. This enabled efficient monitoring and proactive resolution of potential issues.

6.3 Adding Log Monitoring to Analyze Microservice Pod Logs (Using an EFK stack)

6.3.1 Preparation of Minikube and Addition of the Stable Helm Repository

In this step, we prepared the Minikube environment and added the stable Helm repository to enable the installation of necessary Helm charts.

```
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ minikube addons enable default-storageclass
💡 default-storageclass is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
⭐ The 'default-storageclass' addon is enabled
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ minikube addons enable storage-provisioner
💡 storage-provisioner is an addon maintained by minikube. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
    └ Using image gcr.io/k8s-minikube/storage-provisioner:v5
⭐ The 'storage-provisioner' addon is enabled
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ helm repo add stable https://charts.helm.sh/stable
"stable" has been added to your repositories
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ helm repo update
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "fluent" chart repository
...Successfully got an update from the "elastic" chart repository
...Successfully got an update from the "stable" chart repository
Update Complete. *Happy Helming!*
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ █
```

Figure 6.22: Minikube Preparation and Addition of the Stable Helm Repository

Main Objective: Ensure that the Minikube environment is ready to deploy applications with Helm.

6.3.2 Deployment of Elasticsearch with Helm

We installed Elasticsearch using a Helm chart to create an Elasticsearch cluster in the Kubernetes environment.

6.3.2.1 Installation and Verification of Elasticsearch Pods

```
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ helm install elasticsearch stable/elasticsearch
WARNING: This chart is deprecated
NAME: elasticsearch
LAST DEPLOYED: Fri Jan 17 07:23:06 2025
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
This Helm chart is deprecated. Please use https://github.com/elastic/helm-charts/tree/master/elasticsearch instead.

---
The elasticsearch cluster has been installed.

Elasticsearch can be accessed:

* Within your cluster, at the following DNS name at port 9200:
  elasticsearch-client.default.svc

* From outside the cluster, run these commands in the same shell:
  export POD_NAME=$(kubectl get pods --namespace default -l "app=elasticsearch,component=client,release=elasticsearch" -o jsonpath="{.items[0].metadata.name}")
  echo "Visit http://127.0.0.1:9200 to use Elasticsearch"
  kubectl port-forward --namespace default $POD_NAME 9200:9200
```

Figure 6.23: Installation of Elasticsearch with Helm

Once the Elasticsearch pods were deployed, we checked their status to confirm they were running correctly.

NAME	READY	STATUS	RESTARTS	AGE
elasticsearch-client-8755879d9-8wgqq	0/1	PodInitializing	0	10s
elasticsearch-client-8755879d9-hsvtz	0/1	Init:0/1	0	10s
elasticsearch-data-0	0/1	Init:0/2	0	10s
elasticsearch-master-0	0/1	Init:0/2	0	10s

Figure 6.24: Initial Status of Elasticsearch Pods

6.3.2.2 Verification of Elasticsearch Services

We then verified the services exposed by Elasticsearch to confirm that the necessary ports were open.

yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s\$ kubectl get svc						
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	
elasticsearch-client	ClusterIP	10.98.172.134	<none>	9200/TCP	62s	
elasticsearch-discovery	ClusterIP	None	<none>	9300/TCP	62s	
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	5m43s	

Figure 6.25: Verification of Elasticsearch Services

6.3.2.3 Confirmation of Stability of Elasticsearch Pods

After a few minutes, we checked the status of the pods again to ensure they were stable and fully operational.

yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s\$ kubectl get pods				
NAME	READY	STATUS	RESTARTS	AGE
elasticsearch-client-8755879d9-7jft2	1/1	Running	0	8m9s
elasticsearch-client-8755879d9-8s8n6	1/1	Running	0	8m9s
elasticsearch-data-0	1/1	Running	0	8m9s
elasticsearch-data-1	1/1	Running	0	3m12s
elasticsearch-master-0	1/1	Running	0	8m9s
elasticsearch-master-1	1/1	Running	0	3m10s
elasticsearch-master-2	1/1	Running	0	2m52s

Figure 6.26: Status of Elasticsearch Pods After Stabilization

6.3.3 Configuration and Deployment of Fluentd as a DaemonSet

Fluentd was configured as a DaemonSet to collect logs from pods and send them to Elasticsearch.

6.3.3.1 Creation of the Fluentd YAML File

A YAML file was created to define the Fluentd DaemonSet configuration.

```

1  apiVersion: apps/v1
2  kind: DaemonSet
3  metadata:
4    name: fluentd
5    namespace: kube-system
6    labels:
7      k8s-app: fluentd-logging
8      version: v1
9  spec:
10   selector:
11     matchLabels:
12       k8s-app: fluentd-logging
13       version: v1
14   template:
15     metadata:
16       labels:
17         k8s-app: fluentd-logging
18         version: v1
19   spec:
20     tolerations:
21       - key: node-role.kubernetes.io/control-plane
22         effect: NoSchedule
23       - key: node-role.kubernetes.io/master
24         effect: NoSchedule
25     containers:
26       - name: fluentd
27         image: fluent/fluentd-kubernetes-daemonset:v1-debian-elasticsearch
28         env:
29           - name: K8S_NODE_NAME
30             valueFrom:
31               fieldRef:
32                 fieldPath: spec.nodeName
33           - name: FLUENT_ES_HOST
34             value: "elasticsearch-logging"
35           - name: FLUENT_ES_PORT
36             value: "9200"
37           - name: FLUENT_ES_SCHEME
38             value: "http"

```

Figure 6.27: Fluentd YAML Configuration

6.3.3.2 Verification of Fluentd Pods

We deployed Fluentd and verified its pod in the "kube-system" namespace.

```
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ kubectl apply -f fluentd-daemonset-elasticsearch.yaml
daemonset.apps/fluentd created
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ kubectl get pods -n kube-system
NAME                               READY   STATUS            RESTARTS   AGE
coredns-668d6bf9bc-k5jwg          1/1    Running           0          20m
etcd-minikube                      1/1    Running           0          20m
fluentd-4m5vm                     0/1    ContainerCreating   0          28s
```

Figure 6.28: Verification of Fluentd Pod

6.3.4 Deployment of Kibana with a Custom Configuration File

Finally, we deployed Kibana to visualize the logs collected by Fluentd.

6.3.4.1 Configuration of Kibana

The file `kibana-value.yaml` was created to configure Kibana settings.

```
files:
  kibana.yml:
    server.name: kibana
    server.host: "0"
    elasticsearch.hosts: http://elasticsearch-client:9200
service:
  type: ClusterIP
~
```

Figure 6.29: Custom Configuration for Kibana

6.3.4.2 Deployment and Verification of Kibana Pods

We deployed Kibana with Helm and confirmed that the pods were running.

```
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ vi kibana-values.yaml
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ helm install kibana stable/kibana -f kibana-values.yaml
WARNING: This chart is deprecated
NAME: kibana
LAST DEPLOYED: Fri Jan 17 08:54:11 2025
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
THE CHART HAS BEEN DEPRECATED!

Find the new official version @ https://github.com/elastic/helm-charts/tree/master/kibana

To verify that kibana has started, run:

  kubectl --namespace=default get pods -l "app=kibana"

Kibana can be accessed:

  * From outside the cluster, run these commands in the same shell:

    export POD_NAME=$(kubectl get pods --namespace default -l "app=kibana,release=kibana" -o jsonpath=".items[0].metadata.name")
    echo "Visit http://127.0.0.1:5601 to use Kibana"
    kubectl port-forward --namespace default $POD_NAME 5601:5601
```

Figure 6.30: Deployment of Kibana with Helm

NAME	READY	STATUS	RESTARTS	AGE
elasticsearch-client-8755879d9-7jft2	1/1	Running	0	26m
elasticsearch-client-8755879d9-8s8n6	1/1	Running	0	26m
elasticsearch-data-0	1/1	Running	0	26m
elasticsearch-data-1	1/1	Running	0	21m
elasticsearch-master-0	1/1	Running	0	26m
elasticsearch-master-1	1/1	Running	0	21m
elasticsearch-master-2	1/1	Running	0	21m
kibana-5d5d49f79-xxs85	1/1	Running	0	6m26s

Figure 6.31: Verification of Kibana Pods

6.3.4.3 Verification of Service and Access to Kibana

After verifying the service exposed by Kibana, we configured a port-forward to access Kibana from a web browser.

yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s\$ kubectl get svc				
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
elasticsearch-client	ClusterIP	10.109.22.182	<none>	9200/TCP
elasticsearch-discovery	ClusterIP	None	<none>	9300/TCP
kibana	ClusterIP	10.106.246.243	<none>	443/TCP
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP

Figure 6.32: Verification of Kibana Service

```
yassine@yassine-Ubuntu-Data:~/product-microservice-kubernetes/product-k8s$ kubectl port-forward service/kibana 5601:443
Forwarding from 127.0.0.1:5601 -> 5601
Forwarding from [::1]:5601 -> 5601
Handling connection for 5601
```

Figure 6.33: Configuration of Port-Forward to Access Kibana

6.3.4.4 Access to Kibana Dashboard

By accessing the local URL via port-forward, we confirmed that Kibana was functioning correctly.

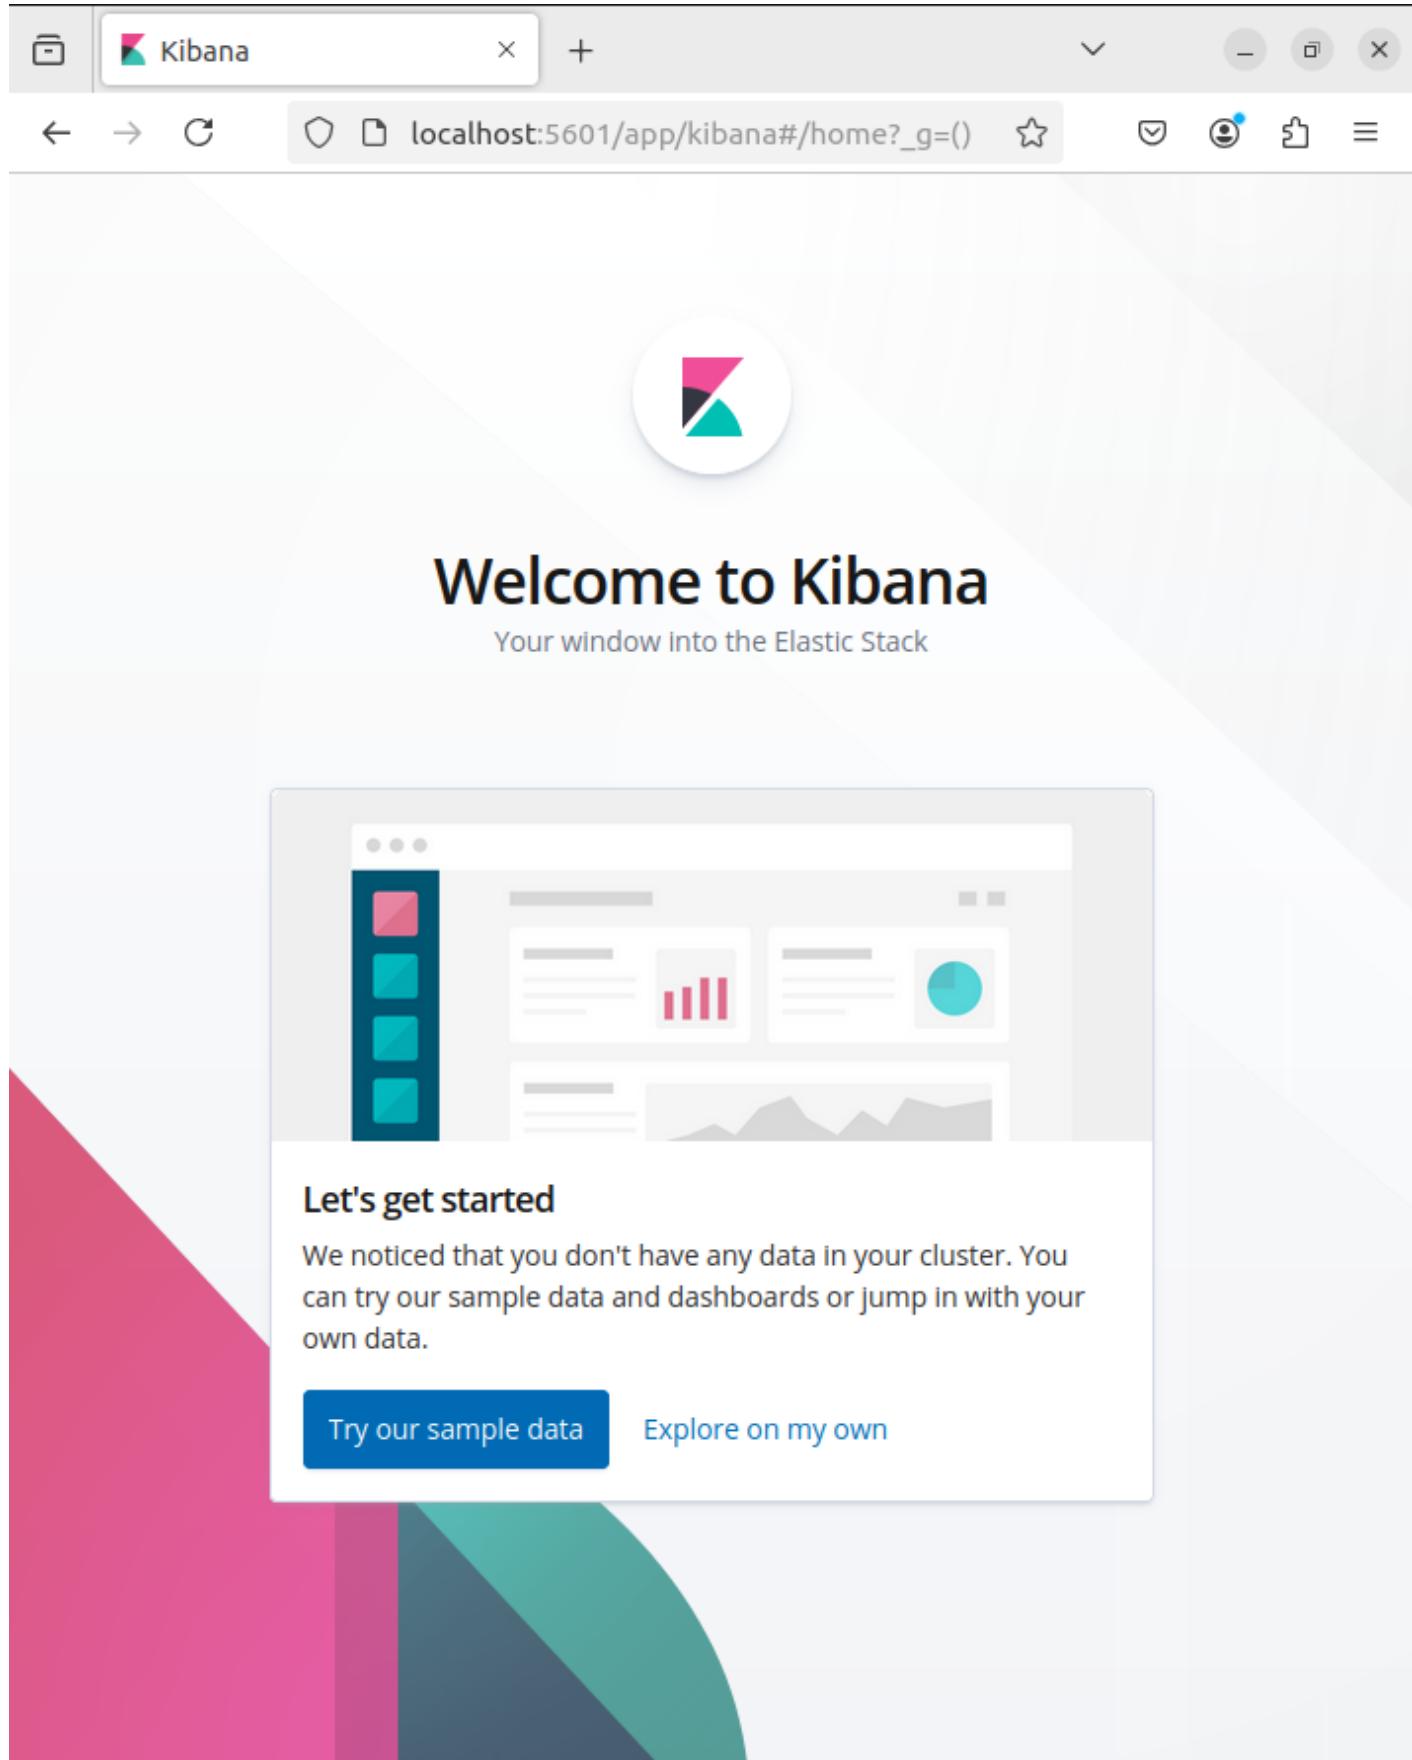


Figure 6.34: Kibana Home Page

6.3.4.5 Log Visualization

Once Kibana was configured, we visualized the logs of the microservice pods to analyze events.

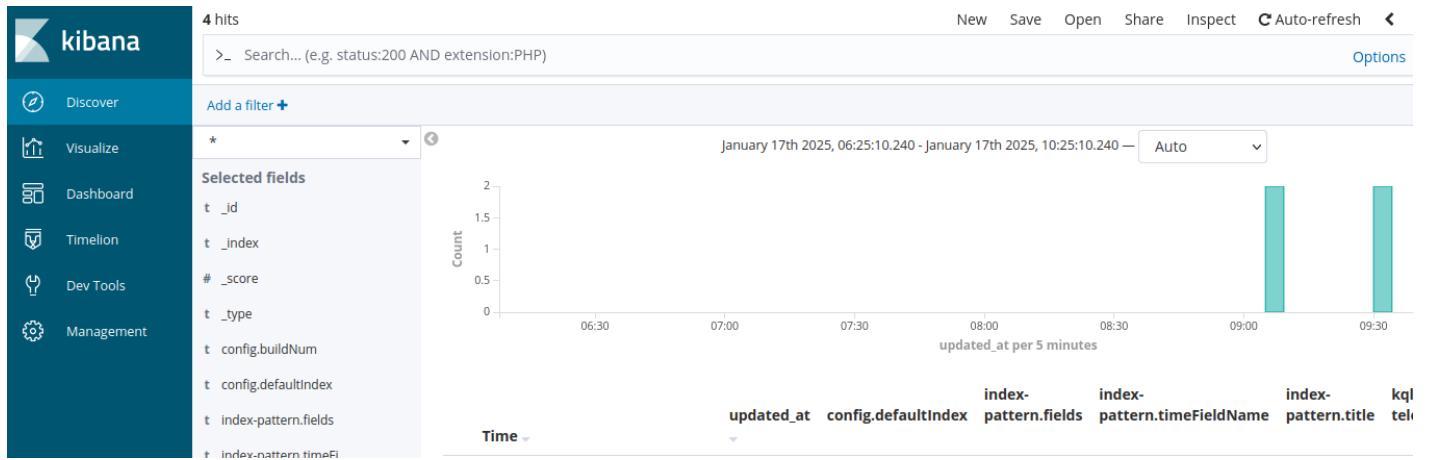


Figure 6.35: Log Visualization in Kibana

6.3.5 Conclusion

With the installation of Elasticsearch, Fluentd, and Kibana, we have set up a complete log monitoring solution to analyze the logs of microservice pods. This architecture ensures better visibility of events and facilitates problem diagnosis in the Kubernetes environment.

General Conclusion and Perspectives

The Proof of Concept (PoC) for automating the build, packaging, and deployment of the Mcommerce application using a DevSecOps approach has demonstrated significant potential for improving KATA's software delivery process. By integrating CI/CD pipelines, containerization, orchestration, and security practices, the project successfully addressed key challenges such as slow delivery cycles, security vulnerabilities, and operational inefficiencies. The results of this PoC highlight the transformative impact of DevSecOps on KATA's ability to deliver high-quality software solutions faster and more securely. While the PoC achieved its objectives, it also revealed areas for improvement, such as the need for more comprehensive testing strategies, further optimization of the CI/CD pipeline. These insights will be invaluable as KATA scales its DevSecOps initiatives across future projects.